

Backend Developer Interview Knowledge Base

AI Interview Coach  RAG Knowledge Base

2026-07-11

Contents

Backend Developer Interview Knowledge Base	2
Backend Architecture	2
REST APIs	3
HTTP Methods	4
Status Codes	5
Authentication	6
Authorization	6
JWT (JSON Web Token)	7
OAuth	8
MVC Architecture	9
Microservices	10
Monolithic Architecture	11
Node.js Overview	11
Express.js Basics	12
Middleware	13
Routing	14
Environment Variables	15
Error Handling	15
Logging	17
Validation	18
Databases	18
SQL	19
NoSQL	20
ORM	21
CRUD	22
API Security	22
Rate Limiting	23
Caching	24

Redis	25
Docker Basics	26
Deployment Concepts	27
Load Balancing	27
Scalability	28
File Upload	28
API Versioning	29
Best Practices (Summary)	30
Performance Optimization	31
Common Interview Mistakes (Cross-Cutting)	32

Backend Developer Interview Knowledge Base

A structured, retrieval-friendly reference for an AI Interview Coach. Each topic is self-contained with definitions, examples, best practices, and common interview mistakes for efficient chunking and semantic search.

Backend Architecture

Definition: Backend architecture describes how server-side components (application logic, databases, APIs, caching, messaging) are organized to handle requests, process data, and return responses to clients.

Core layers in a typical backend: - **Presentation/API layer** - exposes endpoints (REST, GraphQL, gRPC) to clients. - **Business logic layer** - application rules, validation, orchestration. - **Data access layer** - interacts with databases/external services. - **Infrastructure layer** - caching, queues, logging, monitoring.

Client → API Gateway → Service (Controller → Service → Repository) → Database

Best Practices: Separate concerns into layers (controllers handle HTTP, services handle logic, repositories handle data); keep business logic out of controllers; design for statelessness so any server instance can handle any request.

Common Mistakes: Putting database queries directly in controllers (tight coupling, hard to test); mixing business logic across multiple layers inconsistently.

Interview Tip: Be ready to sketch a request's journey from client to database and back, naming each layer's responsibility.

REST APIs

Definition: REST (Representational State Transfer) is an architectural style for designing networked APIs using stateless HTTP requests, resource-based URLs, and standard HTTP methods.

Key principles: statelessness (no session state stored on server between requests), resource-based URLs (/users/123, not /getUser?id=123), uniform interface (standard HTTP verbs), cacheable responses where appropriate.

GET	/api/users	→ list users
GET	/api/users/123	→ get user 123
POST	/api/users	→ create a user
PUT	/api/users/123	→ replace user 123
PATCH	/api/users/123	→ partially update user 123
DELETE	/api/users/123	→ delete user 123

Best Practices: Use nouns (not verbs) in URLs; use plural resource names; nest resources logically (/users/123/orders); return appropriate status codes and consistent JSON structures; support pagination for large collections (?page=2&limit=20).

Common Mistakes: Using verbs in URLs (/getUsers, /createUser) instead of relying on HTTP methods; returning 200 OK for errors; ignoring statelessness by storing session data server-side without a shared store.

Interview Tip: Know the difference between REST and RPC-style APIs, and be ready to explain why REST scales well for stateless, cacheable resource

access.

HTTP Methods

Definition: Standard verbs that indicate the intended action on a resource.

Method	Purpose	Idempotent	Safe	Has Body
GET	Retrieve a resource	Yes	Yes	No
POST	Create a resource	No	No	Yes
PUT	Replace a resource entirely	Yes	No	Yes
PATCH	Partially update a resource	No*	No	Yes
DELETE	Remove a resource	Yes	No	Optional

*Idempotent means repeating the same request produces the same result; PATCH is not guaranteed idempotent depending on implementation.

Example:

```
app.put('/api/users/:id', (req, res) => {  
  // Replaces the entire user resource with req.body  
});  
app.patch('/api/users/:id', (req, res) => {  
  // Updates only the fields provided in req.body  
});
```

Best Practices: Use PUT when sending the complete resource representation; use PATCH for partial updates; never use GET for state-changing operations.

Common Mistakes: Confusing PUT (full replace) with PATCH (partial update); using POST for all operations regardless of semantics.

Status Codes

Definition: Three-digit HTTP response codes indicating the result of a request, grouped by category.

Range	Category	Common Codes
2xx	Success	200 OK, 201 Created, 204 No Content
3xx	Redirection	301 Moved Permanently, 304 Not Modified
4xx	Client Error	400 Bad Request, 401 Unauthorized, 403 Forbidden, 404 Not Found, 409 Conflict, 422 Unprocessable Entity, 429 Too Many Requests
5xx	Server Error	500 Internal Server Error, 502 Bad Gateway, 503 Service Unavailable

Best Practices: Return 201 Created (with a Location header) after successful resource creation; use 204 No Content for successful deletes with no body; use 422 for validation errors versus 400 for malformed requests.

Common Mistakes: Confusing 401 (not authenticated — missing/invalid credentials) with 403 (authenticated but not authorized to access the resource); returning 200 for every response regardless of outcome, forcing clients to parse the body to detect errors.

Interview Tip: Be ready to explain 401 vs 403 with a concrete example — logging in with a bad password is 401; a logged-in user trying to access an admin-only route is 403.

Authentication

Definition: The process of verifying **who** a user is (identity confirmation), typically via credentials (username/password, tokens, biometrics).

Common approaches: session-based (server stores session, client holds a session cookie), token-based (JWT, client holds a signed token), OAuth (delegated authentication via a third party), API keys (for service-to-service auth).

```
// Simplified login flow
app.post('/login', async (req, res) => {
  const user = await User.findOne({ email: req.body.email });
  const valid = await bcrypt.compare(req.body.password, user.passwordHash);
  if (!valid) return res.status(401).json({ error: 'Invalid credentials' });
  const token = jwt.sign({ userId: user.id }, process.env.JWT_SECRET, { ex
  res.json({ token });
});
```

Best Practices: Always hash passwords with a strong algorithm (bcrypt, argon2); never store plaintext passwords; use HTTPS to protect credentials in transit; implement account lockout/rate limiting on login endpoints.

Common Mistakes: Storing passwords in plaintext or with weak hashing (MD5, SHA1 without salting); rolling custom crypto instead of using vetted libraries.

Authorization

Definition: The process of determining **what** an authenticated user is allowed to do (permissions/access control), enforced after authentication.

Common models: Role-Based Access Control (RBAC — permissions tied to roles like admin, user), Attribute-Based Access Control (ABAC — permissions based on resource/user attributes), Access Control Lists (ACL).

```
function requireRole(role) {
  return (req, res, next) => {
    if (req.user.role !== role) {
      return res.status(403).json({ error: 'Forbidden' });
    }
    next();
  };
}
app.delete('/api/users/:id', authenticate, requireRole('admin'), deleteUser)
```

Best Practices: Enforce authorization checks on the server for every protected route — never trust client-side checks alone; apply the principle of least privilege.

Common Mistakes: Confusing authentication with authorization (they are distinct concerns); performing authorization checks only in the UI, allowing direct API calls to bypass them.

Interview Tip: Be ready to state the one-line distinction: *authentication = who you are; authorization = what you can do.*

JWT (JSON Web Token)

Definition: A compact, self-contained, signed token format used to securely transmit claims (user identity, permissions) between parties. Structure: header.payload.signature, Base64Url-encoded.

```
const jwt = require('jsonwebtoken');

// Signing
const token = jwt.sign({ userId: 123, role: 'admin' }, process.env.JWT_SECRET);

// Verifying (middleware)
function authenticate(req, res, next) {
  const token = req.header('Authorization')?.replace('Bearer ', '');
  if (!token) return res.status(401).json({ error: 'Unauthorized' });
  jwt.verify(token, process.env.JWT_SECRET, (err, user) => {
    if (err) return res.status(403).json({ error: 'Forbidden' });
    req.user = user;
    next();
  });
}
```

```
const token = req.headers.authorization?.split(' ')[1];
try {
  req.user = jwt.verify(token, process.env.JWT_SECRET);
  next();
} catch (err) {
  res.status(401).json({ error: 'Invalid or expired token' });
}
}
```

Key properties: stateless (server doesn't store sessions — the token itself carries the claims); signed (not encrypted by default — anyone can decode the payload, but cannot alter it without invalidating the signature).

Best Practices: Keep JWTs short-lived (e.g., 15 min–1 hr) and pair with refresh tokens for longer sessions; store JWTs in HttpOnly cookies (not localStorage) to reduce XSS risk; never put sensitive data (passwords) in the payload since it's readable, not encrypted.

Common Mistakes: Assuming JWTs are encrypted (they're only signed/encoded — Base64 is trivially decodable); storing JWTs in localStorage (vulnerable to XSS token theft); using a weak or hardcoded signing secret.

Interview Tip: Be ready to explain how a server “trusts” a stateless JWT — via signature verification using a secret/public key, without needing a database lookup for every request.

OAuth

Definition: An open standard for **delegated authorization** — allowing a third-party application to access a user's resources on another service without exposing the user's credentials (e.g., “Sign in with Google”).

Key roles: Resource Owner (the user), Client (the app requesting access), Authorization Server (issues tokens), Resource Server (hosts the protected data).

Authorization Code flow (simplified):

1. Client redirects user to Authorization Server login/consent screen.
2. User approves → Authorization Server redirects back with an authorization code.
3. Client exchanges the code (+ client secret) for an access token (server-to-server).
4. Client uses the access token to call the Resource Server's API.

Best Practices: Use the Authorization Code flow with PKCE for public clients (SPAs, mobile apps); never expose client secrets in frontend code; validate the state parameter to prevent CSRF during the OAuth flow.

Common Mistakes: Confusing OAuth (authorization framework) with OpenID Connect (authentication layer built on top of OAuth 2.0, used for identity/login); implementing the Implicit flow (deprecated, less secure) instead of Authorization Code + PKCE.

Interview Tip: Clarify: OAuth answers “can this app access my data?” while authentication (e.g., OpenID Connect) answers “who is this user?” — OAuth alone is not an authentication protocol.

MVC Architecture

Definition: Model-View-Controller is a design pattern separating an application into three interconnected components: **Model** (data/business logic), **View** (presentation, often the client or template), **Controller** (handles requests, coordinates Model and View).

```
// Controller
app.get('/users/:id', async (req, res) => {
  const user = await UserModel.findById(req.params.id); // Model
  res.json(user); // View (JSON representation for an API)
});
```

Best Practices: Keep controllers thin (delegate business logic to services); keep models focused on data structure/persistence; in APIs, the “View” is

often just a serialized JSON response rather than a rendered template.


Common Mistakes: Putting business logic directly in controllers, making them bloated and hard to test; tightly coupling the Model to HTTP-specific concerns (request/response objects).

Interview Tip: In modern API backends, many teams extend MVC to MVCS (Model-View-Controller-Service) — be ready to explain why the service layer was added (to isolate reusable business logic from HTTP handling).

Microservices

Definition: An architectural style where an application is composed of small, independently deployable services, each responsible for a specific business capability, communicating over the network (HTTP, gRPC, message queues).

Characteristics: independent deployment, own database per service (data ownership), communication via APIs/events, technology heterogeneity possible per service.



Best Practices: Design services around bounded contexts (Domain-Driven Design); use asynchronous messaging (queues/event buses) for cross-service communication where possible to reduce coupling; implement centralized logging/tracing (correlation IDs) since requests span multiple services.

Common Mistakes: Creating a “distributed monolith” — services that are technically separate but tightly coupled and must be deployed together; sharing a single database across services (breaks independence); underestimating the operational complexity (network latency, partial failures, distributed transactions).

Interview Tip: Be ready to discuss trade-offs: microservices improve scalability and team autonomy but add complexity in deployment, monitoring, and data consistency (eventual consistency, sagas) compared to a monolith.

Monolithic Architecture

Definition: An architectural style where the entire application (UI, business logic, data access) is built and deployed as a single unit/codebase.

Best Practices: Structure the codebase with clear internal module boundaries (even in a monolith) to ease a future transition to microservices if needed; use a monolith for smaller teams/early-stage products where operational simplicity matters more than independent scaling.

Common Mistakes: Assuming microservices are always superior — for small teams/simple domains, a monolith is often faster to build, test, and deploy, and avoids premature distributed-systems complexity.

Interview Tip: A frequent question: *“When would you choose a monolith over microservices?”* — Answer: smaller teams, early-stage products, simpler operational needs, and when the domain boundaries aren’t yet well understood (avoid premature decomposition).

Node.js Overview

Definition: Node.js is a JavaScript runtime built on Chrome’s V8 engine, enabling server-side JavaScript execution with a non-blocking, event-driven, single-threaded architecture (using an event loop for concurrency).

Key characteristics: single-threaded event loop for I/O-bound concurrency; non-blocking asynchronous I/O (file system, network); npm ecosystem for package management; well-suited for I/O-heavy workloads (APIs, real-time apps) less ideal for CPU-heavy tasks (blocks the event loop).

```
const fs = require('fs');
fs.readFile('data.txt', 'utf8', (err, data) => {
  console.log(data); // non-blocking – other code continues executing meanwhile
});
console.log('This runs before the file read completes');
```

Best Practices: Offload CPU-intensive work to worker threads or separate services to avoid blocking the event loop; use asynchronous APIs (fs.promises, async/await) over synchronous ones (fs.readFileSync) in request-handling code.

Common Mistakes: Using synchronous/blocking calls in request handlers, freezing the entire server for all concurrent requests; assuming Node.js is multi-threaded by default (it uses a single main thread for JS execution, with libuv handling I/O concurrency behind the scenes).

Interview Tip: Be ready to explain why Node.js is well-suited for I/O-bound APIs but a poor fit for CPU-bound tasks (e.g., image processing) without offloading to worker threads or a separate service.

Express.js Basics

Definition: Express.js is a minimal, unopinionated web framework for Node.js used to build APIs and web servers, providing routing, middleware support, and HTTP utility methods.

```
const express = require('express');
const app = express();
app.use(express.json()); // parse JSON request bodies

app.get('/api/health', (req, res) => {
  res.status(200).json({ status: 'ok' });
});
```

```
app.listen(3000, () => console.log('Server running on port 3000'));
```

Best Practices: Organize routes into separate router modules (`express.Router()`) rather than defining all routes in one file; use middleware for cross-cutting concerns (logging, auth, validation).

Common Mistakes: Putting all logic in a single monolithic `app.js` file as the project grows; forgetting `express.json()` middleware, causing `req.body` to be undefined.

Interview Tip: Be ready to build a small CRUD route on the spot — a very common live-coding request.

Middleware

Definition: Functions that execute during the request-response cycle, with access to `req`, `res`, and a `next()` function to pass control to the next middleware/handler. Used for cross-cutting concerns like logging, authentication, and error handling.

```
function logger(req, res, next) {  
  console.log(`${req.method} ${req.url}`);  
  next(); // must call next() to continue the chain  
}  
app.use(logger);  
  
// Error-handling middleware (4 arguments signals Express to treat it as e  
app.use((err, req, res, next) => {  
  console.error(err.stack);  
  res.status(500).json({ error: 'Something went wrong' });  
});
```

Middleware types: application-level (`app.use`), router-level, error-handling (4 params), built-in (`express.json()`), third-party (`cors`, `helmet`, `morgan`).

Best Practices: Keep middleware focused on a single concern; place error-handling middleware **last**, after all routes; always call `next()` (or send a response) to avoid hanging requests.

Common Mistakes: Forgetting to call `next()`, causing the request to hang indefinitely; placing error-handling middleware before regular routes (it won't catch their errors).

Interview Tip: Be ready to explain middleware execution order — Express processes middleware top-to-bottom in the order they're registered.

Routing

Definition: The mechanism that maps HTTP request URLs and methods to specific handler functions.

```
const router = express.Router();

router.get('/', getAllUsers);
router.get('/:id', getUserById);
router.post('/', createUser);
router.put('/:id', updateUser);
router.delete('/:id', deleteUser);

app.use('/api/users', router); // mount router with a base path
```

Best Practices: Group related routes into dedicated router modules by resource; use route parameters (`:id`) for resource identifiers and query parameters (`?sort=asc`) for filtering/sorting/pagination; version routes when introducing breaking changes (`/api/v1/users`).

Common Mistakes: Defining overly generic catch-all routes that unintentionally shadow more specific ones (route order matters — more specific routes should be defined before generic/wildcard ones).

Environment Variables

Definition: Configuration values (API keys, database URLs, secrets, ports) stored outside the codebase, loaded at runtime, allowing different configuration per environment (development, staging, production) without code changes.

```
// .env file (never committed to version control)  
// DATABASE_URL=postgres://user:pass@localhost:5432/mydb  
// JWT_SECRET=supersecretvalue  
  
require('dotenv').config();  
const dbUrl = process.env.DATABASE_URL;  
const port = process.env.PORT || 3000;
```

Best Practices: Never commit .env files or secrets to version control (add to .gitignore); provide a .env.example template documenting required variables without real values; validate required environment variables at startup and fail fast if missing.

Common Mistakes: Hardcoding secrets/config directly in source code; committing .env files with real credentials to a public repository (a very common real-world security incident).

Interview Tip: Be ready to explain how environment-specific configuration supports the “twelve-factor app” principle of strict separation between config and code.

Error Handling

Definition: The practice of anticipating, catching, and gracefully responding to runtime errors (validation failures, database errors, unexpected exceptions) instead of letting them crash the server or expose internals.

```

// Centralized error-handling middleware
class AppError extends Error {
  constructor(message, statusCode) {
    super(message);
    this.statusCode = statusCode;
  }
}

app.get('/api/users/:id', async (req, res, next) => {
  try {
    const user = await User.findById(req.params.id);
    if (!user) throw new AppError('User not found', 404);
    res.json(user);
  } catch (err) {
    next(err); // delegate to error-handling middleware
  }
});

app.use((err, req, res, next) => {
  const status = err.statusCode || 500;
  res.status(status).json({ error: err.message || 'Internal Server Error'
});

```

Best Practices: Use a centralized error-handling middleware instead of duplicating try/catch response logic in every route; distinguish operational errors (expected, e.g., validation) from programmer errors (bugs); never leak stack traces or internal details to clients in production.

Common Mistakes: Swallowing errors silently (empty catch blocks); returning inconsistent error response shapes across different endpoints; crashing the entire process on an unhandled promise rejection without graceful recovery/logging.

Interview Tip: Be ready to explain the difference between handling errors gracefully (returning a proper 4xx/5xx response) versus letting the pro-

cess crash — and how to catch unhandled rejections at the process level (`process.on('unhandledRejection', ...)`) as a safety net.

Logging

Definition: Recording application events (requests, errors, state changes) to help with debugging, monitoring, and auditing in production.

```
const winston = require('winston');
const logger = winston.createLogger({
  level: 'info',
  format: winston.format.json(),
  transports: [new winston.transports.Console(), new winston.transports.Fi
});

logger.info('Server started', { port: 3000 });
logger.error('Database connection failed', { error: err.message });
```

Best Practices: Use structured logging (JSON) for easier querying/aggregation in log management tools (ELK, Datadog); include contextual data (request ID, user ID) for traceability; use appropriate log levels (debug, info, warn, error); never log sensitive data (passwords, tokens, PII).

Common Mistakes: Using `console.log` everywhere in production code instead of a proper logging library with levels/transports; logging sensitive information in plaintext.

Interview Tip: Be ready to explain the value of correlation/request IDs for tracing a single request across multiple microservices in distributed systems.

Validation

Definition: Verifying that incoming data (request bodies, query params) meets expected structure, types, and constraints before processing, to prevent invalid data and security vulnerabilities.

```
const { body, validationResult } = require('express-validator');

app.post('/api/users',
  body('email').isEmail().withMessage('Invalid email'),
  body('password').isLength({ min: 8 }).withMessage('Password too short'),
  (req, res) => {
    const errors = validationResult(req);
    if (!errors.isEmpty()) {
      return res.status(422).json({ errors: errors.array() });
    }
    // proceed with valid data
  }
);
```

Best Practices: Validate on the server even if client-side validation exists (never trust the client); use schema validation libraries (Joi, Zod, express-validator) rather than ad-hoc manual checks; validate early, before touching the database.

Common Mistakes: Relying solely on client-side validation (trivially bypassed via direct API calls); inconsistent or missing validation across endpoints, leading to malformed data reaching the database.

Databases

Definition: Systems for storing, organizing, and retrieving structured or unstructured data, broadly categorized as relational (SQL) or non-relational (NoSQL).

Best Practices: Choose the database type based on data shape and access patterns, not familiarity alone; use connection pooling to avoid exhausting database connections under load; index frequently queried columns.

Common Mistakes: Choosing a database technology purely by popularity without evaluating consistency, scalability, and query pattern needs.

Interview Tip: Be ready to justify a database choice for a given scenario (e.g., “would you use SQL or NoSQL for an e-commerce order system?”).

SQL

Definition: Structured Query Language — used with relational databases (PostgreSQL, MySQL) that store data in structured tables with predefined schemas and relationships (foreign keys), and support ACID transactions.

```
SELECT u.name, o.total
FROM users u
JOIN orders o ON u.id = o.user_id
WHERE o.total > 100
ORDER BY o.total DESC
LIMIT 10;
```

```
INSERT INTO users (name, email) VALUES ('Alice', 'alice@example.com');
UPDATE users SET email = 'new@example.com' WHERE id = 1;
DELETE FROM users WHERE id = 1;
```

ACID properties: Atomicity (all-or-nothing transactions), Consistency (valid state transitions), Isolation (concurrent transactions don’t interfere), Durability (committed data persists).

Best Practices: Use parameterized queries/prepared statements to prevent SQL injection; normalize schemas to reduce redundancy, but denormalize selectively for read-heavy performance needs; index columns used in WHERE/JOIN/ORDER BY.

Common Mistakes: String-concatenating user input directly into SQL queries (critical SQL injection vulnerability); over-normalizing to the point of requiring excessive joins for common queries.

Interview Tip: Be ready to explain SQL injection with an example and how parameterized queries prevent it:

```
// Vulnerable:
db.query(`SELECT * FROM users WHERE email = '${email}'`);
// Safe:
db.query('SELECT * FROM users WHERE email = ?', [email]);
```

NoSQL

Definition: Non-relational databases that store data in flexible formats (documents, key-value, wide-column, graph) rather than fixed tables — optimized for scalability and schema flexibility over strict relational consistency.

Types: Document (MongoDB — JSON-like documents), Key-Value (Redis, DynamoDB), Wide-Column (Cassandra), Graph (Neo4j).

```
// MongoDB example
db.users.insertOne({ name: 'Alice', email: 'alice@example.com', tags: ['ad
db.users.find({ tags: 'admin' }));
```

Best Practices: Use NoSQL for flexible/evolving schemas, high write throughput, or horizontal scaling needs; design document structure around access patterns (denormalize/embed related data when read together frequently).

Common Mistakes: Using NoSQL for highly relational data requiring complex joins and strong transactional consistency across entities; ignoring eventual consistency implications in distributed NoSQL systems.

Interview Tip: Be ready to compare SQL vs NoSQL trade-offs: SQL offers strong consistency and relational integrity; NoSQL offers flexible schemas

and easier horizontal scaling — the choice depends on data shape and consistency requirements (CAP theorem trade-offs).

ORM

Definition: Object-Relational Mapping — a library that lets developers interact with a relational database using object-oriented code instead of writing raw SQL (e.g., Sequelize, TypeORM, Prisma for Node.js; Hibernate for Java).

```
// Prisma example
const user = await prisma.user.create({
  data: { name: 'Alice', email: 'alice@example.com' }
});
const users = await prisma.user.findMany({ where: { active: true } });
```

Best Practices: Use an ORM for productivity and protection against SQL injection (parameterization is handled internally); drop down to raw queries for complex, performance-critical queries the ORM can't express efficiently.

Common Mistakes: The “N+1 query problem” — fetching a list of records and then issuing a separate query per record for related data instead of using eager loading/joins:

```
// N+1 problem: one query per user to fetch their orders
users.forEach(async user => { user.orders = await Order.find({ userId: user.id }); });
// Fix: eager load with a single join query
const users = await User.findAll({ include: Order });
```

Interview Tip: The N+1 query problem is a very common ORM interview question — be ready to explain both the problem and the fix (eager loading/joins/batching).

CRUD

Definition: Create, Read, Update, Delete — the four basic operations for persistent data, typically mapped to HTTP methods (POST, GET, PUT/PATCH, DELETE) in REST APIs.

```
router.post('/users', createUser);    // Create
router.get('/users/:id', getUser);    // Read
router.put('/users/:id', updateUser); // Update
router.delete('/users/:id', deleteUser); // Delete
```

Best Practices: Ensure each CRUD operation returns the appropriate status code (201 for create, 200/204 for update/delete); wrap multi-step writes in database transactions to maintain consistency.

Common Mistakes: Not using transactions for operations that modify multiple related records, risking partial/inconsistent writes if one step fails.

API Security

Definition: Practices and mechanisms to protect APIs from unauthorized access, data breaches, and abuse.

Key measures: HTTPS everywhere (encrypt data in transit); input validation/sanitization (prevent injection attacks); authentication + authorization on every protected route; rate limiting (prevent abuse/DoS); security headers via middleware (helmet in Express); CORS configuration (restrict allowed origins); secrets management (environment variables, vaults — never hard-coded).

```
const helmet = require('helmet');
const cors = require('cors');

app.use(helmet()); // sets secure HTTP headers
app.use(cors({ origin: 'https://trusted-frontend.com', credentials: true }));
```

Best Practices: Follow the principle of least privilege for all credentials/roles; keep dependencies updated to patch known vulnerabilities; log and monitor for suspicious activity; sanitize all user input server-side regardless of client-side checks.

Common Mistakes: Trusting client-side validation alone; exposing detailed error messages/stack traces in production responses; using wildcard CORS (*) with credentialed requests.

Interview Tip: Be ready to name the OWASP Top 10 categories relevant to APIs (e.g., broken authentication, injection, broken access control, security misconfiguration) even at a high level.

Rate Limiting

Definition: Restricting the number of requests a client can make to an API within a given time window, to prevent abuse, protect infrastructure, and ensure fair usage.

```
const rateLimit = require('express-rate-limit');

const limiter = rateLimit({
  windowMs: 15 * 60 * 1000, // 15 minutes
  max: 100,                  // limit each IP to 100 requests per window
  message: 'Too many requests, please try again later.'
});
app.use('/api/', limiter);
```

Common algorithms: Fixed window, Sliding window, Token bucket, Leaky bucket.

Best Practices: Apply stricter limits on sensitive endpoints (login, password reset) to mitigate brute-force attacks; return 429 Too Many Requests with a Retry-After header when limits are exceeded; use a shared store (Redis) for rate limiting across multiple server instances.

Common Mistakes: Implementing rate limiting with in-memory storage in a multi-instance deployment (each instance tracks limits independently, making the limit ineffective) instead of a shared/distributed store.

Caching

Definition: Temporarily storing frequently accessed data in a fast-access layer (memory) to reduce latency and database/computation load on repeated requests.

Levels: client-side (browser cache, Cache-Control headers), CDN caching (static assets), application-level caching (in-memory or Redis), database query caching.

```
app.get('/api/products/:id', async (req, res) => {
  const cacheKey = `product:${req.params.id}`;
  const cached = await redisClient.get(cacheKey);
  if (cached) return res.json(JSON.parse(cached));

  const product = await Product.findById(req.params.id);
  await redisClient.setEx(cacheKey, 3600, JSON.stringify(product)); // cac
  res.json(product);
});
```

Cache invalidation strategies: TTL (time-based expiry), write-through (update cache on write), cache-aside/lazy loading (populate cache on read miss, as shown above).

Best Practices: Cache read-heavy, infrequently-changing data; always set a reasonable TTL to avoid serving stale data indefinitely; invalidate/update cache explicitly on writes when strong freshness matters.

Common Mistakes: Caching highly volatile data without proper invalidation, serving stale results; the classic quote applies here — “There are only two hard things in Computer Science: cache invalidation and naming things.”

Interview Tip: Be ready to explain cache-aside vs write-through strategies and when each is appropriate.

Redis

Definition: An in-memory, key-value data store used for caching, session storage, rate limiting, pub/sub messaging, and as a lightweight message broker — valued for its speed (sub-millisecond latency).

```
const redis = require('redis');
const client = redis.createClient();

await client.set('user:123', JSON.stringify({ name: 'Alice' })), { EX: 3600 };
const value = await client.get('user:123');

await client.lPush('queue:jobs', JSON.stringify({ task: 'sendEmail' })); //
await client.expire('user:123', 60); // set/update TTL
```

Common use cases: caching, session store (shared across multiple server instances), rate limiting counters, real-time leaderboards (sorted sets), pub/sub for real-time features.

Best Practices: Set expiration (EX) on cache keys to avoid unbounded memory growth; use Redis as a shared session store in horizontally-scaled Node.js apps instead of in-memory sessions.

Common Mistakes: Using Redis as a primary/durable data store for critical data without proper persistence configuration (RDB/AOF); storing unbounded, ever-growing data without eviction policies.

Docker Basics

Definition: Docker packages an application and its dependencies into a portable, isolated **container** that runs consistently across environments (development, staging, production).

```
FROM node:20-alpine
WORKDIR /app
COPY package*.json ./
RUN npm install --production
COPY . .
EXPOSE 3000
CMD ["node", "server.js"]
```

```
docker build -t my-app .
docker run -p 3000:3000 -e NODE_ENV=production my-app
```

Key concepts: **Image** (immutable snapshot/blueprint), **Container** (running instance of an image), **Dockerfile** (instructions to build an image), **docker-compose** (define/run multi-container applications, e.g., app + database + Redis together).

Best Practices: Use small base images (alpine variants) to reduce image size; use multi-stage builds to keep production images lean (exclude dev dependencies/build tools); never bake secrets into images — pass them via environment variables at runtime.

Common Mistakes: Including node_modules/build artifacts in version control and copying them into the image instead of installing fresh; running containers as the root user unnecessarily (security risk).

Interview Tip: Be ready to explain the difference between an image and a container (class vs instance analogy) and why containers ensure “it works on my machine” consistency across environments.

Deployment Concepts

Definition: The process and strategies for releasing application code to production environments reliably and with minimal downtime.

Common strategies: **Blue-Green deployment** (two identical environments, switch traffic instantly between them), **Rolling deployment** (gradually replace old instances with new ones), **Canary deployment** (release to a small subset of users first, monitor, then roll out fully).

CI/CD: Continuous Integration (automatically build/test code on every commit) + Continuous Deployment/Delivery (automatically deploy passing builds to production/staging).

Best Practices: Automate deployments via CI/CD pipelines to reduce human error; use health checks so load balancers/orchestrators only route traffic to healthy instances; always have a rollback plan.

Common Mistakes: Deploying manually without automated tests/health checks, increasing the risk of shipping broken code to production; no rollback strategy when a deployment fails.

Load Balancing

Definition: Distributing incoming network traffic across multiple server instances to improve availability, reliability, and horizontal scalability.

Common algorithms: Round Robin (requests distributed sequentially), Least Connections (routes to the server with fewest active connections), IP Hash (routes based on client IP for session stickiness).

Client → Load Balancer → [Server 1] [Server 2] [Server 3]

Best Practices: Design backend services to be stateless so any instance can handle any request (store session state in a shared store like Redis, not in-memory); use health checks to automatically remove unhealthy instances from rotation.

Common Mistakes: Relying on server-side in-memory session state in a load-balanced environment, causing inconsistent behavior when a user's requests hit different servers ("session affinity" issues).

Interview Tip: Be ready to explain why statelessness is essential for effective load balancing and horizontal scaling.

Scalability

Definition: A system's ability to handle increased load (more users, more data, more requests) by adding resources.

Vertical scaling (scale up): adding more power (CPU/RAM) to an existing server — simple but has a hard ceiling and a single point of failure. **Horizontal scaling (scale out):** adding more server instances behind a load balancer — more complex but virtually unlimited and more resilient.

Best Practices: Design for horizontal scalability from the start where growth is expected (stateless services, externalized session/cache state); scale the database independently via read replicas, sharding, or caching layers, since it's often the bottleneck before application servers are.

Common Mistakes: Assuming scaling application servers alone solves performance issues while ignoring the database as the actual bottleneck; premature scaling investment before establishing real load requirements.

Interview Tip: Be ready to discuss how you'd scale a specific system component (e.g., "how would you scale a read-heavy API?") — typical answers include caching, read replicas, and CDN usage.

File Upload

Definition: Handling binary file data (images, documents) sent from clients to the server, typically via multipart/form-data encoding.

```

const multer = require('multer');
const upload = multer({
  dest: 'uploads/',
  limits: { fileSize: 5 * 1024 * 1024 }, // 5MB limit
  fileFilter: (req, file, cb) => {
    const allowed = ['image/png', 'image/jpeg'];
    cb(null, allowed.includes(file.mimetype));
  }
});

app.post('/api/upload', upload.single('file'), (req, res) => {
  res.json({ filename: req.file.filename });
});

```

Best Practices: Validate file type and size on the server (never trust the client’s declared MIME type alone — verify actual content when security matters); store uploaded files in object storage (S3, Cloud Storage) rather than the application server’s local disk in production (ephemeral/non-scalable); generate unique filenames to avoid collisions/overwrites.

Common Mistakes: Not limiting file size, allowing denial-of-service via huge uploads; trusting the client-supplied file extension/MIME type without validation, allowing malicious file uploads (e.g., disguised executables).

Interview Tip: Be ready to explain why uploaded files shouldn’t be stored on a single server’s local disk in a horizontally-scaled, load-balanced architecture (files wouldn’t be accessible from other instances) — object storage or a shared file service solves this.

API Versioning

Definition: Strategies for evolving an API over time without breaking existing client integrations.

Common approaches:

- **URI versioning:** /api/v1/users, /api/v2/users — simplest, most visible.
- **Header versioning:** Accept: application/vnd.myapi.v2+json — cleaner URLs, less discoverable.
- **Query parameter versioning:** /api/users?version=2 — simple but easy to overlook.

```
app.use('/api/v1/users', usersV1Router);  
app.use('/api/v2/users', usersV2Router);
```

Best Practices: Version APIs from the start of any public-facing API to avoid painful retrofits later; maintain backward compatibility within a version (only introduce breaking changes in a new version); communicate deprecation timelines clearly to API consumers.

Common Mistakes: Making breaking changes to an existing API version without warning, breaking client integrations; over-versioning trivial, non-breaking changes unnecessarily.

Interview Tip: Be ready to explain what counts as a “breaking change” (removing/renaming fields, changing data types, altering required parameters) versus a safe additive change (adding new optional fields).

Best Practices (Summary)

- Design stateless, horizontally scalable services; externalize session state (Redis) when load balancing.
- Validate and sanitize all input server-side, regardless of client-side checks.
- Use parameterized queries/ORMs to prevent SQL injection.
- Hash passwords with bcrypt/argon2; never store plaintext credentials.
- Return consistent, meaningful HTTP status codes and error response shapes.
- Use environment variables for configuration/secrets; never commit them to version control.
- Implement centralized error handling and structured logging.

- Apply rate limiting and security headers (helmet) to all public APIs.
 - Cache read-heavy data with clear invalidation strategy (TTL or explicit invalidation).
 - Write automated tests and use CI/CD pipelines for reliable, repeatable deployments.
-

Performance Optimization

Definition: Techniques to reduce latency, increase throughput, and use server resources efficiently under load.

Key techniques: - **Caching** at multiple layers (CDN, application, database query cache). - **Database indexing** on frequently queried/joined columns. - **Connection pooling** to reuse database connections instead of opening new ones per request. - **Pagination** for large result sets instead of returning entire datasets. - **Asynchronous/non-blocking I/O** to maximize throughput per server instance. - **Compression** (gzip/Brotli) for HTTP responses. - **Load testing** (e.g., k6, Apache JMeter) to identify bottlenecks before they hit production.

```
// Example: pagination to avoid loading entire collections
app.get('/api/products', async (req, res) => {
  const page = parseInt(req.query.page) || 1;
  const limit = parseInt(req.query.limit) || 20;
  const products = await Product.find()
    .skip((page - 1) * limit)
    .limit(limit);
  res.json(products);
});
```

Best Practices: Profile and measure before optimizing (avoid premature optimization); optimize the biggest bottleneck first (often the database, not the application code); use indexes deliberately (too many indexes slow down writes).

Common Mistakes: Optimizing code without profiling first, wasting effort on parts that aren't actual bottlenecks; missing database indexes on frequently filtered/joined columns, causing full table scans; N+1 query patterns silently degrading performance under load.

Interview Tip: Be ready to walk through how you would diagnose a “slow API endpoint” — check database query performance/indexes first, then look at N+1 queries, then caching opportunities, then application-level bottlenecks.

Common Interview Mistakes (Cross-Cutting)

1. **Confusing authentication and authorization** — unable to clearly state the distinction with an example.
2. **Confusing 401 and 403** status codes.
3. **Not knowing the difference between PUT and PATCH** (full replace vs partial update).
4. **Storing JWTs insecurely** (localStorage instead of HttpOnly cookies) or assuming JWTs are encrypted.
5. **Not validating input server-side**, relying only on frontend validation.
6. **SQL injection vulnerabilities** from string-concatenated queries instead of parameterized queries.
7. **Ignoring the N+1 query problem** when discussing ORM usage.
8. **Assuming microservices are always the better choice** without discussing trade-offs (operational complexity, data consistency).
9. **Not knowing how to scale a stateful service** — missing the connection between statelessness and horizontal scaling/load balancing.
10. **Hardcoding secrets/config** instead of using environment variables.
11. **Not having a clear caching invalidation strategy** when discussing performance improvements.
12. **Overlooking rate limiting and basic API security measures** (CORS, helmet, HTTPS) when designing a “production-ready” API.

Interview Tip: When asked an open-ended system design question, struc-

ture your answer: clarify requirements → outline high-level architecture → discuss data model → discuss scaling/security/failure handling — this structured approach signals strong engineering communication skills.

End of Knowledge Base.