

Data Analyst Interview Coach — Knowledge Base

A Comprehensive Reference for RAG-Based Interview Preparation

Interview Coach Knowledge Base

2026

Contents

Introduction to This Knowledge Base	3
Part 1: The Data Analysis Process	3
1.1 What Is the Data Analysis Process?	3
1.2 Defining the Business Problem	4
1.3 Data Collection Sources	4
1.4 Structuring an Analysis Request	4
1.5 Data Types and Data Grain	5
1.6 Communicating with Stakeholders	5
Part 2: Excel for Data Analysis	5
2.1 Excel Basics	5
2.2 Advanced Excel Functions	6
2.3 Pivot Tables	7
2.4 VLOOKUP	7
2.5 XLOOKUP	7
2.6 Power Query	8
Part 3: SQL for Data Analysis	8
3.1 SELECT and Filtering with WHERE	8
3.2 GROUP BY and Aggregation	8
3.3 HAVING vs. WHERE	9
3.3b Common Table Expressions (CTEs)	9
3.3c String, Date, and CASE Functions	9
3.3d UNION vs. UNION ALL	10
3.4 JOINS	10
3.5 Subqueries	11
3.6 Window Functions	11
3.7 Views	12
3.8 Indexes	12
3.9 Normalization	12
3.10 Primary Keys, Foreign Keys, and Star Schemas	13
3.11 Query Performance Basics	13

Part 4: Python for Data Analysis	14
4.1 NumPy	14
4.2 Pandas	14
4.3 Data Cleaning in Python	15
4.4 Exploratory Data Analysis (EDA)	16
4.5 Matplotlib	17
Part 5: BI Tools — Power BI and Tableau	17
5.1 Power BI	17
5.2 Tableau	18
5.3 Data Modeling in BI Tools	19
5.4 Live vs. Extract / Import vs. DirectQuery	19
Part 6: Statistics for Data Analysts	20
6.1 Measures of Central Tendency	20
6.2 Standard Deviation and Variance	20
6.3 Probability Fundamentals	20
6.3b Sampling and the Central Limit Theorem	21
6.4 Hypothesis Testing	22
6.5 A/B Testing	22
Part 7: Business Metrics, Dashboards, and Visualization	23
7.1 Business Metrics and KPIs	23
7.1b Funnel, Retention, and Cohort Analysis	24
7.2 Dashboards	24
7.3 Data Visualization Principles	25
7.4 Storytelling with Data	25
Part 8: Common Interview Mistakes and Best Practices	25
8.1 Common Mistakes Candidates Make	25
8.2 Best Practices for Interviews	26
8.2b The STAR Framework for Behavioral Questions	27
8.3 Sample Behavioral Interview Questions	27
8.4 Sample Technical Interview Prompts	28
Appendix A: Quick-Reference Cheat Sheets	28
A.1 SQL Clause Execution Order	28
A.2 Excel Lookup Function Comparison	28
A.3 SQL JOIN Types at a Glance	28
A.4 Central Tendency and Spread Cheat Sheet	28
A.5 Statistical Test Selection Guide	29
A.6 Chart Type Selection Guide	29
A.7 Pandas vs. SQL Equivalents	29
Appendix B: Glossary of Key Terms	30
Appendix C: Rapid-Fire Interview Question Bank	31
Excel Keyboard Shortcuts Reference	34

Appendix D: Worked SQL Practice Problems	34
Appendix E: Worked Python/Pandas Practice Problems	35
Appendix F: Sample Case Study Walkthrough	36
Appendix G: Portfolio Project Ideas for Data Analysts	37
Appendix H: 30-Day Interview Preparation Checklist	38
Appendix I: Skill Expectations by Experience Level	38
Closing Notes for the Interview Coach	39

Introduction to This Knowledge Base

This document is a structured knowledge base for a Data Analyst interview preparation assistant. Each section is self-contained and written to be retrieved independently by a semantic search system. Sections include concise definitions, interview-style questions and answers, and practical examples. Use this document to answer candidate questions on data analysis concepts, tools, and interview strategy.

Part 1: The Data Analysis Process

1.1 What Is the Data Analysis Process?

The data analysis process is the structured workflow analysts follow to turn raw data into actionable insight. The standard stages are:

1. **Define the Problem** — Clarify the business question, stakeholders, and success metrics.
2. **Data Collection** — Gather data from databases, APIs, files, or third-party sources.
3. **Data Cleaning** — Handle missing values, duplicates, outliers, and formatting issues.
4. **Exploratory Data Analysis (EDA)** — Understand distributions, relationships, and patterns.
5. **Analysis / Modeling** — Apply statistical methods, aggregations, or machine learning.
6. **Visualization** — Represent findings in charts, dashboards, or reports.

7. **Communication / Storytelling** — Present insights and recommendations to stakeholders.
8. **Iteration** — Refine based on feedback and monitor outcomes over time.

Interview Q&A: - *Q: Walk me through your data analysis process.* A: Strong answers mention starting with the business question, understanding the data source, cleaning the data, exploring it, running an analysis appropriate to the question, and closing the loop by presenting a clear recommendation tied to business impact — not just the numbers. - *Q: How do you know when a dataset is “ready” for analysis?* A: When missing values, duplicates, and type mismatches are resolved, key fields are validated against expected ranges, and the grain of the data (what one row represents) is confirmed.

1.2 Defining the Business Problem

Before touching data, analysts should identify: the decision the analysis will inform, the audience, the timeline, and the metric that defines success. A common interview trap is jumping straight into SQL/Excel without first restating the business question in your own words.

Best practice: Always restate the ask (“So we want to understand why signups dropped in March, correct?”) before beginning technical work — this is a top signal interviewers look for.

1.3 Data Collection Sources

Common sources: relational databases (SQL), data warehouses (Snowflake, BigQuery, Redshift), flat files (CSV, Excel), APIs, web scraping, and third-party survey tools. Analysts should be able to describe the data lineage — where data originates and how it flows into the tables they query.

1.4 Structuring an Analysis Request

When a stakeholder brings a vague request (“Can you look into why engagement is down?”), analysts should convert it into a testable question with a defined metric, timeframe, and segment. For example: “Did 7-day active users decline week-over-week for mobile users in the US between weeks 10 and 12, and if so, in which funnel stage?”

Interview Q&A: - *Q: A stakeholder asks you to “just pull the numbers.” How do you respond?* A: Ask what decision the numbers will inform, over what time period, and

what “success” looks like, so the output is directly usable rather than a raw data dump requiring rework. - *Q: How do you scope a request when you have limited time?* A: Identify the minimum viable analysis that still answers the core business question, communicate what’s in and out of scope up front, and flag follow-up areas for a later pass.

1.5 Data Types and Data Grain

Before analysis, confirm the **grain** of each table (what a single row represents — e.g., one row per order, per customer, per day). Mismatched grain is a common cause of incorrect JOINS and inflated aggregates. Also confirm data types: numeric, string, date/time, boolean, and categorical — implicit type mismatches (e.g., a date stored as text) silently break comparisons and sorting.

Interview Q&A: - *Q: You JOIN two tables and your row count doubles. What likely happened?* A: A grain mismatch — one table likely has multiple rows per key (e.g., multiple line items per order), causing a fan-out during the JOIN that inflates any subsequent SUM or COUNT.

1.6 Communicating with Stakeholders

Strong analysts translate between technical and business language. Practices include: confirming requirements in writing, setting expectations on turnaround time, presenting a headline insight before the details, and proactively flagging data limitations (e.g., “this excludes returns” or “sample size is small for this segment”).

Part 2: Excel for Data Analysis

2.1 Excel Basics

Core Excel skills expected of analysts: cell referencing (relative vs. absolute, e.g., \$A\$1), basic formulas (SUM, AVERAGE, COUNT, COUNTIF, SUMIF), conditional formatting, data validation, sorting/filtering, and freeze panes for large datasets.

Interview Q&A: - *Q: What’s the difference between a relative and absolute reference?* A: A relative reference (A1) shifts when copied to other cells; an absolute reference (\$A\$1) stays fixed. Mixed references (\$A1 or A\$1) lock only the row or column. - *Q: How do COUNTIF and SUMIF differ from COUNTIFS and SUMIFS?* A: The “S” versions support multiple criteria across multiple ranges; the base versions support only

one condition. - *Q: What's the difference between Sort and Filter in Excel?* A: Sort permanently reorders rows based on one or more columns; Filter temporarily hides rows that don't meet criteria without changing their order.

Useful Everyday Formulas

Formula	Purpose
=SUMIFS(sum_range, crit_range1, crit1, crit_range2, crit2)	Sum with multiple conditions
=COUNTIFS(range1, crit1, range2, crit2)	Count with multiple conditions
=AVERAGEIFS(avg_range, crit_range1, crit1)	Average with conditions
=TRIM(A1)	Remove extra spaces from text
=LEFT(A1,3) / =RIGHT(A1,3) / =MID(A1,2,4)	Extract substrings
=TEXT(A1, "yyyy-mm-dd")	Format a value as text
=UNIQUE(range)	Return distinct values (modern Excel)
=SORT(range)	Sort a range dynamically (modern Excel)

2.2 Advanced Excel Functions

Key advanced functions: INDEX/MATCH, IF/IFS, SUMPRODUCT, TEXT, CONCATENATE/TEXTJOIN, array formulas, and IFERROR for error handling. Advanced analysts also use named ranges and structured table references (Table1[Column]).

Interview Q&A: - *Q: Why might INDEX/MATCH be preferred over VLOOKUP?* A: INDEX/MATCH can look left, isn't broken by inserted columns, and is generally faster on large datasets. - *Q: How do you handle errors gracefully in a formula?* A: Wrap the formula in IFERROR(formula, fallback_value) so a lookup or division failure returns a clean value (like 0 or "N/A") instead of an error code. - *Q: What's the difference between IF and nested IFS?* A: A single IF handles one condition with a true/false branch; IFS (or nested IFs) evaluates multiple sequential conditions, similar to a SQL CASE WHEN statement.

Data Validation and Error Handling

Data Validation restricts what can be entered into a cell (e.g., dropdown lists, number ranges, date ranges), reducing data entry errors at the source — an important quality-

control habit for analysts who build shared input templates.

Common error types: #N/A (lookup not found), #REF! (invalid cell reference, often from deleted rows/columns), #DIV/0! (division by zero), #VALUE! (wrong data type in a formula).

2.3 Pivot Tables

Pivot Tables summarize large datasets by dragging fields into Rows, Columns, Values, and Filters. They support aggregation (sum, count, average, min/max), grouping (by date, by range), calculated fields, and slicers for interactive filtering.

Interview Q&A: - *Q: How would you find total sales by region and month using a Pivot Table?* A: Put Region in Rows, Month in Columns, and Sales in Values (set to Sum). Add a slicer for further interactivity. - *Q: What's a calculated field in a Pivot Table?* A: A custom formula built from existing fields (e.g., Profit = Revenue – Cost) that appears as a new value in the pivot.

2.4 VLOOKUP

VLOOKUP(lookup_value, table_array, col_index_num, [range_lookup]) searches for a value in the first column of a range and returns a value from a specified column in the same row.

=VLOOKUP(A2, Sheet2!A:D, 3, FALSE)

Limitations: VLOOKUP only searches left-to-right (lookup column must be leftmost), breaks if columns are inserted/deleted, and FALSE must be specified for exact match (default is approximate, which is a common bug source).

Interview Q&A: - *Q: What does the fourth argument in VLOOKUP do?* A: It sets exact (FALSE/0) vs. approximate (TRUE/1) match. Analysts should almost always use FALSE unless doing range-based lookups like tax brackets.

2.5 XLOOKUP

XLOOKUP(lookup_value, lookup_array, return_array, [if_not_found], [match_mode], [search_mode]) is the modern replacement for VLOOKUP/HLOOKUP. It searches in any direction, defaults to exact match, supports a built-in “not found” message, and can return multiple columns at once.

=XLOOKUP(A2, Sheet2!A:A, Sheet2!C:C, "Not Found")

Interview Q&A: - *Q: Why is XLOOKUP considered better than VLOOKUP?* A: It defaults to exact match (safer), can search left or right, handles errors gracefully with `if_not_found`, and isn't affected by inserted/deleted columns since it references arrays, not column index numbers.

2.6 Power Query

Power Query is Excel's (and Power BI's) ETL tool for importing, transforming, and combining data from multiple sources without formulas. Common operations: removing duplicates, splitting columns, unpivoting, merging queries (joins), appending queries (unions), and changing data types. Steps are recorded in the "Applied Steps" pane using the M language, enabling repeatable, refreshable transformations.

Interview Q&A: - *Q: When would you use Power Query instead of formulas?* A: When transformations need to be repeatable/refreshable, when combining multiple files or sources, or when the transformation logic (unpivot, merge, type conversion) is complex enough that formulas become unmanageable. - *Q: What's the difference between Merge and Append in Power Query?* A: Merge joins two tables side-by-side based on a key (like a SQL JOIN); Append stacks tables vertically (like a SQL UNION).

Part 3: SQL for Data Analysis

3.1 SELECT and Filtering with WHERE

SELECT retrieves columns from a table; WHERE filters rows before aggregation.

```
SELECT customer_id, order_date, amount
FROM orders
WHERE amount > 100 AND order_date >= '2026-01-01';
```

Interview Q&A: - *Q: What's the execution order of a SQL query?* A: Logically: FROM → JOIN → WHERE → GROUP BY → HAVING → SELECT → ORDER BY → LIMIT. This is a very common interview question. - *Q: Can you use a column alias defined in SELECT inside the WHERE clause?* A: No — WHERE executes before SELECT, so aliases aren't yet defined. You'd need to repeat the expression or use a subquery/CTE.

3.2 GROUP BY and Aggregation

GROUP BY groups rows sharing a value so aggregate functions (SUM, COUNT, AVG, MIN, MAX) can be applied per group.

```
SELECT region, COUNT(*) AS num_orders, SUM(amount) AS total_sales
FROM orders
GROUP BY region;
```

Rule: Every non-aggregated column in SELECT must appear in GROUP BY.

3.3 HAVING vs. WHERE

HAVING filters groups after aggregation, while WHERE filters rows before aggregation.

```
SELECT region, SUM(amount) AS total_sales
FROM orders
GROUP BY region
HAVING SUM(amount) > 10000;
```

Interview Q&A: - Q: *Why can't you filter on an aggregate using WHERE?* A: Because WHERE is evaluated before aggregation happens, so aggregate values don't exist yet at that stage.

3.3b Common Table Expressions (CTEs)

A CTE, defined with WITH, creates a temporary named result set that can be referenced within a single query — improving readability over nested subqueries and allowing reuse of the same logic multiple times.

```
WITH regional_sales AS (
    SELECT region, SUM(amount) AS total_sales
    FROM orders
    GROUP BY region
)
SELECT region, total_sales
FROM regional_sales
WHERE total_sales > 10000;
```

Interview Q&A: - Q: *Why use a CTE instead of a subquery?* A: CTEs are more readable, can be referenced multiple times in the same query, and support recursion (recursive CTEs), which regular subqueries cannot do.

3.3c String, Date, and CASE Functions

Common string functions: UPPER(), LOWER(), TRIM(), SUBSTRING(), CONCAT(), LENGTH(). Common date functions: DATEADD(), DATEDIFF(), DATE_TRUNC(), EXTRACT().

CASE WHEN implements conditional logic in SQL, similar to Excel's IF.

```
SELECT customer_id,  
       CASE  
         WHEN total_spent > 1000 THEN 'High Value'  
         WHEN total_spent > 100 THEN 'Mid Value'  
         ELSE 'Low Value'  
       END AS customer_tier  
FROM customer_totals;
```

3.3d UNION vs. UNION ALL

UNION combines result sets from two queries and removes duplicates; UNION ALL combines them and keeps all rows, including duplicates, making it faster since it skips deduplication.

Interview Q&A: - Q: When should you use UNION ALL instead of UNION? A: When you know there are no duplicates to worry about, or duplicates are meaningful (e.g., combining transaction logs) — UNION ALL avoids the performance cost of a distinct/sort operation.

3.4 JOINS

- **INNER JOIN** — returns rows with matches in both tables.
- **LEFT JOIN** — returns all rows from the left table, with NULLs where no match exists on the right.
- **RIGHT JOIN** — mirror of LEFT JOIN.
- **FULL OUTER JOIN** — returns all rows from both tables, matched where possible.
- **SELF JOIN** — a table joined to itself, often for hierarchical data (e.g., employee-manager).
- **CROSS JOIN** — Cartesian product of two tables.

```
SELECT o.order_id, c.customer_name  
FROM orders o  
LEFT JOIN customers c ON o.customer_id = c.customer_id;
```

Interview Q&A: - Q: How do you find customers with no orders? A: LEFT JOIN customers to orders and filter WHERE orders.order_id IS NULL. - Q: What's the difference between INNER JOIN and LEFT JOIN in terms of row count? A: INNER JOIN can only return rows with matches (equal or fewer rows than either table); LEFT JOIN always

retains every row of the left table regardless of match. - *Q: What happens if you put a filter on the right table in the WHERE clause vs. the ON clause of a LEFT JOIN?*

A: A condition in ON is applied during the join (preserving unmatched left rows with NULLs); the same condition in WHERE is applied after the join and will drop unmatched left rows, effectively turning the LEFT JOIN into an INNER JOIN. - *Q: What is a many-to-many join and why is it risky?* A: When neither side of the join key is unique, each matching combination is returned, which can cause row counts (and downstream sums) to explode unexpectedly — always check key uniqueness before joining.

3.5 Subqueries

A subquery is a query nested inside another query, used in SELECT, FROM, or WHERE clauses.

```
SELECT customer_id
FROM orders
WHERE amount > (SELECT AVG(amount) FROM orders);
```

Correlated vs. non-correlated: A correlated subquery references the outer query and runs once per outer row (slower); a non-correlated subquery runs independently once.

Interview Q&A: - *Q: When would you use a subquery instead of a JOIN?* A: When you need an aggregate comparison (e.g., “above average”) or a single scalar value, subqueries are often clearer. JOINS are typically preferred for combining row-level data due to performance.

3.6 Window Functions

Window functions perform calculations across a set of rows related to the current row without collapsing them, using OVER().

```
SELECT employee_id, department, salary,
       RANK() OVER (PARTITION BY department ORDER BY salary DESC) AS salary_rank,
       AVG(salary) OVER (PARTITION BY department) AS dept_avg
FROM employees;
```

Common window functions: ROW_NUMBER(), RANK(), DENSE_RANK(), LAG(), LEAD(), SUM() OVER(), AVG() OVER().

Interview Q&A: - *Q: What's the difference between RANK, DENSE_RANK, and ROW_NUMBER?* A: ROW_NUMBER() assigns unique sequential numbers even for ties; RANK() gives ties the same rank but skips subsequent numbers (1,1,3); DENSE_RANK()

gives ties the same rank without skipping (1,1,2). - *Q: How would you calculate month-over-month growth in SQL?* A: Use `LAG(revenue) OVER (ORDER BY month)` to get the prior month's value, then compute `(current - previous) / previous`.

3.7 Views

A view is a saved, virtual table defined by a query. It doesn't store data itself (unless materialized) — it re-runs the underlying query each time it's referenced.

```
CREATE VIEW high_value_customers AS
SELECT customer_id, SUM(amount) AS total_spent
FROM orders
GROUP BY customer_id
HAVING SUM(amount) > 5000;
```

Interview Q&A: - *Q: Why use a view instead of a table?* A: Views simplify complex/repeated queries, provide a consistent interface for analysts, and can restrict access to sensitive columns without duplicating data. - *Q: What's a materialized view?* A: A view whose results are physically stored and periodically refreshed, trading storage/staleness for faster read performance.

3.8 Indexes

An index is a data structure (commonly a B-tree) that speeds up row lookups on specified columns, at the cost of extra storage and slower writes (inserts/updates must maintain the index).

Interview Q&A: - *Q: When would you add an index?* A: On columns frequently used in WHERE, JOIN, or ORDER BY clauses, especially on large tables where full table scans are expensive. - *Q: What's a downside of over-indexing?* A: Slower write performance and increased storage, since every insert/update/delete must also update each index.

3.9 Normalization

Normalization organizes relational data to reduce redundancy and improve integrity, typically through Normal Forms:

- **1NF** — Atomic values, no repeating groups.
- **2NF** — 1NF + no partial dependency on a composite key.
- **3NF** — 2NF + no transitive dependency (non-key columns depend only on the key).

Denormalization intentionally introduces redundancy (e.g., in a data warehouse) to optimize read performance for analytics.

Interview Q&A: - *Q: Why might a data warehouse be denormalized while a transactional (OLTP) database is normalized?* A: OLTP systems prioritize data integrity and fast writes; data warehouses (OLAP) prioritize fast, simple reads for analytics, so star/snowflake schemas trade some redundancy for query performance and simplicity.

3.10 Primary Keys, Foreign Keys, and Star Schemas

A **primary key** uniquely identifies each row in a table; a **foreign key** is a column referencing a primary key in another table, enforcing referential integrity. A **star schema** is a common data-warehouse design with a central **fact table** (transactions/events, numeric measures) surrounded by **dimension tables** (descriptive attributes like customer, product, date).

Interview Q&A: - *Q: What's the difference between a fact table and a dimension table?* A: A fact table stores quantitative, transactional data (e.g., sales amount, quantity) at a defined grain; dimension tables store descriptive context (customer name, product category) used to filter and group facts. - *Q: Why use a surrogate key instead of a natural key?* A: Surrogate keys (auto-incrementing IDs) are stable even if business attributes change (e.g., a customer's email), and they simplify joins across a warehouse's star schema.

3.11 Query Performance Basics

Analysts are often expected to reason about query performance even without deep DBA knowledge. Key habits: filter early (WHERE before large JOINS where possible), avoid SELECT * on wide tables, use indexed columns in JOIN/WHERE clauses, and check the query execution plan (EXPLAIN) when a query is unexpectedly slow.

Interview Q&A: - *Q: Your dashboard query used to run in 2 seconds and now takes 2 minutes. How do you investigate?* A: Check whether data volume grew significantly, whether an index was dropped or is no longer used, review the execution plan for full table scans, and confirm no new JOINS or subqueries were added that create row explosion.

Part 4: Python for Data Analysis

4.1 NumPy

NumPy provides fast, memory-efficient n-dimensional arrays and vectorized operations, forming the numerical foundation for pandas and most data science libraries.

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5])
mean = arr.mean()
std = arr.std()
filtered = arr[arr > 2]
```

Interview Q&A: - *Q: Why is NumPy faster than native Python lists for numerical operations?* A: NumPy arrays are stored contiguously in memory as a single data type and use vectorized C-level operations, avoiding Python's per-element loop overhead.

4.2 Pandas

Pandas is the core library for tabular data manipulation, built on NumPy. Core objects: Series (1D) and DataFrame (2D).

```
import pandas as pd
df = pd.read_csv("sales.csv")
df.head()
df.groupby("region")["sales"].sum()
df[df["amount"] > 100]
df.merge(other_df, on="customer_id", how="left")
df.pivot_table(values="sales", index="region", columns="month", aggfunc="sum")
```

Interview Q&A: - *Q: What's the difference between .loc and .iloc?* A: .loc selects by label (row/column names); .iloc selects by integer position. - *Q: How would you find duplicate rows in a DataFrame?* A: df.duplicated() returns a boolean mask; df.drop_duplicates() removes them. - *Q: How does groupby in pandas relate to GROUP BY in SQL?* A: Conceptually identical — both split data into groups and apply an aggregate function per group. - *Q: How do you apply multiple aggregations at once in pandas?* A: Use .agg(), e.g. df.groupby("region")["sales"].agg(["sum", "mean", "count"]), or pass a dictionary to apply different functions per column. - *Q: What's the difference between .apply() and a vectorized operation?* A: Vectorized operations (e.g., df["a"] + df["b"]) run in optimized C code and are much faster; .apply() runs a Python function row-by-row (or column-by-column) and is slower — use it only when a vectorized alternative doesn't exist. - *Q: How*

do the merge “how” options (inner, left, right, outer) compare to SQL joins? A: `df.merge(other, on="key", how="inner"/"left"/"right"/"outer")` maps directly to SQL’s INNER/LEFT/RIGHT/FULL OUTER JOIN.

Additional Pandas Patterns

Reshaping

```
df_wide = df.pivot(index="date", columns="category", values="sales")
df_long = df.melt(id_vars=["date"], var_name="category", value_name="sales")
```

Datetime handling

```
df["month"] = df["date"].dt.to_period("M")
df["day_of_week"] = df["date"].dt.day_name()
```

Window-style calculations

```
df["rolling_avg"] = df["sales"].rolling(window=7).mean()
df["rank"] = df.groupby("region")["sales"].rank(ascending=False)
```

Concatenating and combining

```
combined = pd.concat([df1, df2], axis=0) # stack rows (like UNION ALL)
```

4.3 Data Cleaning in Python

Common data cleaning tasks: handling missing values (`df.isnull()`, `df.fillna()`, `df.dropna()`), removing duplicates, fixing data types (`astype`), standardizing text (`str.lower()`, `str.strip()`), handling outliers (IQR or z-score methods), and renaming/reshaping columns.

```
df["date"] = pd.to_datetime(df["date"])
df["amount"] = df["amount"].fillna(df["amount"].median())
Q1, Q3 = df["amount"].quantile([0.25, 0.75])
IQR = Q3 - Q1
df = df[(df["amount"] >= Q1 - 1.5*IQR) & (df["amount"] <= Q3 + 1.5*IQR)]
```

Interview Q&A: - Q: How do you decide whether to drop or impute missing values? A: Depends on the proportion missing, whether the missingness is random (MCAR/MAR) or systematic (MNAR), and the analysis’s sensitivity. Small random gaps in non-critical fields can often be dropped; important fields might be imputed with mean/median/mode or a model-based estimate. - Q: What’s the IQR method for outlier detection? A: Outliers are values below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$,

where $IQR = Q3 - Q1$.

Data Cleaning Checklist

A practical checklist analysts can run through on any new dataset:

- ☐ Confirm the grain (what one row represents) and check for unexpected duplicate keys.
- ☐ Check data types for every column; fix columns where numbers/dates are stored as text.
- ☐ Quantify missing values per column and decide: drop, impute, or flag.
- ☐ Standardize categorical text values (casing, whitespace, spelling variants like “USA” vs. “United States”).
- ☐ Check numeric ranges for impossible values (e.g., negative ages, future birth dates).
- ☐ Identify and investigate outliers rather than automatically removing them.
- ☐ Validate foreign-key-style relationships (e.g., every `customer_id` in orders exists in customers).
- ☐ Re-check row counts before and after each transformation step to catch silent data loss or duplication.

Interview Q&A: - *Q: How do you handle a categorical field with inconsistent labels (e.g., “NY,” “New York,” “ny”)?* A: Standardize via case normalization, trimming whitespace, and a mapping/lookup table for known synonyms, ideally validated against a canonical reference list. - *Q: What’s the difference between a missing value that’s random vs. systematic?* A: Random missingness (MCAR/MAR) doesn’t relate meaningfully to other variables and can often be safely imputed; systematic missingness (MNAR) correlates with the value itself (e.g., high earners skipping an income field) and naive imputation can bias results.

4.4 Exploratory Data Analysis (EDA)

EDA is the process of summarizing a dataset’s main characteristics, often visually, before formal modeling. Typical steps: check shape and data types (`df.info()`, `df.describe()`), examine distributions (histograms, boxplots), check correlations (heatmaps), and identify missing values/outliers.

Interview Q&A: - *Q: What’s the first thing you do when you open a new dataset?* A: Check its shape, column types, and summary statistics (`df.shape`, `df.info()`, `df.describe()`), then look for missing values and obvious anomalies before deeper analysis. - *Q: How do you check for correlations between numeric variables?* A:

`df.corr()` produces a correlation matrix, often visualized with a heatmap (via `seaborn` or `matplotlib`) to quickly spot strong positive/negative relationships. - Q: *What's the risk of skipping EDA and jumping straight to modeling or reporting?* A: You may miss data quality issues (outliers, skew, leakage, incorrect types) that silently bias results, and you may choose an inappropriate method for the data's actual distribution.

4.5 Matplotlib

Matplotlib is Python's foundational plotting library, giving fine-grained control over charts.

```
import matplotlib.pyplot as plt
plt.figure(figsize=(8,5))
plt.plot(df["month"], df["sales"])
plt.title("Monthly Sales")
plt.xlabel("Month")
plt.ylabel("Sales")
plt.show()
```

Common chart types: line (`plt.plot`), bar (`plt.bar`), scatter (`plt.scatter`), histogram (`plt.hist`), boxplot (`plt.boxplot`).

Interview Q&A: - Q: *When would you use a boxplot vs. a histogram?* A: A histogram shows the full shape/frequency of a distribution; a boxplot is more compact and better for comparing distributions across categories or spotting outliers quickly.

Part 5: BI Tools — Power BI and Tableau

5.1 Power BI

Power BI is Microsoft's business intelligence tool for building interactive dashboards. Key components: Power Query (data transformation), Data Model (relationships between tables), DAX (Data Analysis Expressions for calculated columns/measures), and Report view (visualizations).

Key DAX concepts: - **Calculated Column** — computed row-by-row, stored in the table. - **Measure** — computed dynamically based on filter context (e.g., `Total Sales = SUM(Sales[Amount])`).

Total Sales = SUM(Sales[Amount])

YoY Growth = DIVIDE([Total Sales] - [Total Sales PY], [Total Sales PY])

Interview Q&A: - *Q: What's the difference between a calculated column and a measure in Power BI?* A: A calculated column is evaluated per row and stored in the model (uses more memory); a measure is evaluated dynamically at query time based on the current filter/visual context (more efficient for aggregations). - *Q: What is filter context in DAX?* A: The set of filters (from slicers, rows/columns in a visual, or explicit filters) applied when a measure is evaluated, which determines the subset of data the calculation runs on.

5.2 Tableau

Tableau is a leading data visualization tool emphasizing drag-and-drop chart building. Key concepts: Dimensions (categorical fields) vs. Measures (numeric fields), Live vs. Extract connections, calculated fields, LOD (Level of Detail) expressions (FIXED, INCLUDE, EXCLUDE), and dashboards combining multiple worksheets.

Interview Q&A: - *Q: What's the difference between a Dimension and a Measure in Tableau?* A: Dimensions are qualitative fields used to slice data (e.g., Region, Category); Measures are quantitative, aggregatable fields (e.g., Sales, Profit). - *Q: What is a LOD expression used for?* A: To control the level of aggregation independent of the visualization's granularity — e.g., {FIXED [Customer ID] : SUM([Sales])} computes total sales per customer regardless of what's on the view. - *Q: Power BI vs. Tableau — how would you compare them?* A: Power BI integrates tightly with the Microsoft ecosystem and is generally more cost-effective with strong data modeling (DAX); Tableau is often praised for more flexible, polished visual design and strong ad hoc exploration. Many teams choose based on existing tech stack and licensing.

DAX Time Intelligence Functions

Power BI's DAX includes built-in functions for common date-based comparisons, which rely on a proper Date table marked as the model's date dimension:

Sales PY = CALCULATE([Total Sales], SAMEPERIODLASTYEAR('Date'[Date]))

Sales YTD = TOTALYTD([Total Sales], 'Date'[Date])

Sales MTD = TOTALMTD([Total Sales], 'Date'[Date])

Interview Q&A: - *Q: Why do time intelligence functions require a dedicated Date table?* A: They rely on a continuous, unbroken date range marked as the model's official date dimension to correctly calculate periods like year-to-date or same-period-

last-year; using a fact table's date column directly can produce incorrect results if dates are missing or not continuous. - *Q: What does the CALCULATE function do in DAX?* A: It evaluates an expression within a modified filter context, making it the most important and versatile function in DAX — most time intelligence and conditional aggregation patterns are built on top of it.

5.3 Data Modeling in BI Tools

Both tools rely on relationships between tables (similar to SQL joins) rather than flattening everything into one table. Star schemas are preferred for performance and clarity. Analysts should understand cardinality (one-to-many vs. many-to-many relationships) and filter direction, since incorrect relationships are a leading cause of wrong numbers in dashboards.

Interview Q&A: - *Q: Why is it best practice to model data as a star schema in Power BI/Tableau rather than one big flat table?* A: Star schemas reduce data redundancy, improve refresh/query performance, and make relationships explicit, whereas a single flat table can cause double-counting issues in aggregations when dimensions repeat. - *Q: What is a blended data source in Tableau, and when would you use it?* A: Data blending combines data from different connections (e.g., a SQL database and an Excel file) at an aggregate level using a common field, useful when a true join isn't possible across systems.

5.4 Live vs. Extract / Import vs. DirectQuery

Tableau's **Live** connection queries the source in real time; an **Extract** caches a snapshot for faster performance. Power BI's **Import** mode loads data into memory (fast, but requires refresh); **DirectQuery** queries the source live (always current, but slower and more limited).

Interview Q&A: - *Q: When would you choose a live connection over an extract?* A: When data changes frequently and near-real-time accuracy is critical (e.g., operational dashboards); extracts are preferred when performance matters more than second-by-second freshness (e.g., daily executive reporting).

Part 6: Statistics for Data Analysts

6.1 Measures of Central Tendency

- **Mean** — the arithmetic average; sensitive to outliers.
- **Median** — the middle value when sorted; robust to outliers.
- **Mode** — the most frequently occurring value; useful for categorical data.

Interview Q&A: - *Q: When would you use median instead of mean?* A: When the data is skewed or contains outliers (e.g., income, home prices), since the median better represents a “typical” value.

6.2 Standard Deviation and Variance

Variance measures average squared deviation from the mean; **standard deviation** is its square root, expressed in the same units as the data, making it more interpretable.

$$\text{Variance} = \Sigma(x - \text{mean})^2 / n$$

$$\text{Std Dev} = \sqrt{\text{Variance}}$$

Interview Q&A: - *Q: Why do we square the deviations when calculating variance?* A: To eliminate negative values that would otherwise cancel out and to penalize larger deviations more heavily. - *Q: What does a low standard deviation tell you?* A: Data points are clustered closely around the mean, indicating low variability/consistency.

6.3 Probability Fundamentals

Core concepts: probability rules (addition, multiplication), conditional probability $P(A|B) = P(A \cap B) / P(B)$, independence, and common distributions (normal, binomial, Poisson).

Interview Q&A: - *Q: What's the difference between independent and mutually exclusive events?* A: Independent events don't affect each other's probability; mutually exclusive events cannot occur together (and are therefore actually dependent — if one happens, the other's probability becomes zero).

Common Distributions

- **Normal distribution** — symmetric, bell-shaped; described by mean and standard deviation; underlies many statistical tests via the Central Limit Theorem.
- **Binomial distribution** — number of successes in a fixed number of independent yes/no trials (e.g., conversion counts).

- **Poisson distribution** — number of events in a fixed interval of time/space, useful for modeling rare events (e.g., website errors per hour).

Correlation vs. Causation

Correlation measures how two variables move together (commonly via Pearson's r , ranging from -1 to 1); it does **not** imply that one causes the other. Confounding variables, reverse causation, and coincidence can all produce correlation without causation.

Interview Q&A: - *Q: Give an example of correlation without causation.* A: Ice cream sales and drowning incidents both rise in summer — they're correlated through a confounding variable (hot weather), not because one causes the other. - *Q: How would you try to establish causation from observational data?* A: Control for confounders (regression adjustment), look for natural experiments/instrumental variables, or ideally run a randomized controlled experiment (A/B test) where feasible.

6.3b Sampling and the Central Limit Theorem

Analysts rarely have full population data, so they work with samples. Key concepts: **sampling bias** (a non-representative sample skews results), **random sampling** (every unit has an equal chance of selection), and the **Central Limit Theorem (CLT)** — the sampling distribution of the mean approaches a normal distribution as sample size grows, regardless of the population's original distribution. This justifies using normal-based confidence intervals and t-tests even on non-normal data, given a large enough sample.

Confidence Intervals

A confidence interval gives a range of plausible values for a population parameter, with a stated confidence level (e.g., 95%). A 95% CI means that if we repeated the sampling process many times, about 95% of the constructed intervals would contain the true population value.

Interview Q&A: - *Q: What does "95% confidence" actually mean?* A: It describes the long-run reliability of the estimation method, not the probability that this specific interval contains the true parameter — a common misconception worth clarifying in interviews.

6.4 Hypothesis Testing

Hypothesis testing evaluates whether observed data provides enough evidence to reject a **null hypothesis (H0)** in favor of an **alternative hypothesis (H1)**.

Key concepts: - **p-value** — probability of observing results at least as extreme as the sample, assuming H0 is true. - **Significance level (α)** — the threshold (commonly 0.05) below which we reject H0. - **Type I error** — false positive (rejecting a true H0). - **Type II error** — false negative (failing to reject a false H0). - **Confidence interval** — a range likely to contain the true population parameter.

Common tests: t-test (compare means), chi-square test (categorical association), ANOVA (compare 3+ group means).

Interview Q&A: - *Q: What does a p-value of 0.03 mean?* A: If the null hypothesis were true, there'd be a 3% chance of observing a result this extreme or more by random chance — below the typical 0.05 threshold, so we'd reject H0. - *Q: What's the difference between Type I and Type II error?* A: Type I is rejecting a true null hypothesis (false alarm); Type II is failing to reject a false null hypothesis (missed effect). - *Q: What's the difference between a one-tailed and two-tailed test?* A: A one-tailed test checks for an effect in one specific direction (e.g., "is the new version better?"); a two-tailed test checks for a difference in either direction (e.g., "is the new version different, better or worse?"). Two-tailed tests are the more conservative default unless there's a strong reason to test only one direction.

Worked Example: A website redesign is tested against the old design. Old design: 400 conversions out of 10,000 visitors (4.0%). New design: 460 conversions out of 10,000 visitors (4.6%). A two-proportion z-test would compare these rates, accounting for sample size, to determine whether a 0.6 percentage-point lift is statistically significant or could plausibly be due to random variation. Interviewers often want candidates to reason about *why* larger samples make small differences easier to detect (narrower standard error), even without computing the exact z-statistic by hand.

6.5 A/B Testing

A/B testing compares two versions (control vs. treatment) to determine which performs better on a defined metric, using randomized assignment and statistical significance testing.

Process: Define hypothesis and metric → determine sample size/power → randomly assign users → run test for a sufficient duration → analyze results with a significance test → check for practical significance, not just statistical significance.

Interview Q&A: - *Q: How do you determine when an A/B test has run long enough?*

A: By pre-calculating the required sample size based on baseline conversion rate, minimum detectable effect, and desired power/significance — then running until that sample size is reached, avoiding “peeking” that inflates false positive risk. - *Q: What’s the difference between statistical significance and practical significance?* A: Statistical significance means the observed difference is unlikely due to chance; practical significance means the effect size is large enough to matter for the business. A result can be statistically significant but too small to justify action. - *Q: What’s a common pitfall in A/B testing?* A: Stopping the test early based on interim results (“peeking”), which inflates the false positive rate; also failing to account for novelty effects or seasonality.

Part 7: Business Metrics, Dashboards, and Visualization

7.1 Business Metrics and KPIs

A **KPI (Key Performance Indicator)** is a measurable value showing how effectively a business is achieving key objectives. Good KPIs are specific, measurable, tied to a business goal, and trackable over time.

Common examples: Revenue, Conversion Rate, Customer Acquisition Cost (CAC), Customer Lifetime Value (CLV), Churn Rate, Monthly Active Users (MAU), Net Promoter Score (NPS).

Interview Q&A: - *Q: How do you decide which metrics matter for a given business question?* A: Trace back from the business objective to the specific decision it informs, then choose metrics that are directly actionable, not just easy to measure (“vanity metrics”).

- *Q: What’s the difference between a leading and a lagging indicator?*

A: Leading indicators predict future performance (e.g., signups this week); lagging indicators reflect past outcomes (e.g., last quarter’s revenue).

SaaS and Subscription Metrics

- **MRR (Monthly Recurring Revenue)** — predictable monthly revenue from subscriptions.
- **ARR (Annual Recurring Revenue)** — $\text{MRR} \times 12$; used for annual planning.
- **CAC (Customer Acquisition Cost)** — $\text{total sales \& marketing spend} \div \text{number of new customers acquired}$.

- **LTV:CAC Ratio** — compares customer lifetime value to acquisition cost; a ratio above ~3:1 is often considered healthy.
- **Churn Rate** — percentage of customers or revenue lost in a period.
- **Net Revenue Retention (NRR)** — revenue retained from existing customers including upgrades/downgrades/churn, often expressed as a percentage (above 100% indicates expansion outweighs churn).

Interview Q&A: - *Q: Why might a company with high revenue growth still be in trouble?* A: If CAC is rising faster than LTV, or churn is high and masked by strong new acquisition, the growth may be unsustainable — a classic scenario analysts are asked to diagnose using the underlying unit economics rather than top-line revenue alone. - *Q: What's the difference between gross churn and net churn?* A: Gross churn only counts revenue lost from cancellations/downgrades; net churn nets that loss against expansion revenue (upsells) from the same existing customer base, and can even be negative if expansion outpaces losses.

7.1b Funnel, Retention, and Cohort Analysis

A **funnel** breaks a process into sequential steps (e.g., Visit → Sign Up → Purchase) and measures the conversion rate between each step, helping identify where users drop off. **Retention** measures the percentage of users who return after a given period (Day 1, Day 7, Day 30 retention are common). **Cohort analysis** groups users by a shared starting point (e.g., signup week) and tracks their behavior over time, isolating trends that a single blended metric would hide.

Interview Q&A: - *Q: Why use cohort analysis instead of just looking at overall retention?* A: Overall retention can mask trends — e.g., improving retention among new cohorts might be hidden by a large, aging cohort with naturally declining engagement. Cohorts isolate behavior by starting group so trends are comparable. - *Q: How would you identify the biggest drop-off point in a signup funnel?* A: Calculate the conversion rate between each consecutive step and compare against historical benchmarks or segments; the step with the largest percentage-point drop relative to expectation is the priority to investigate.

7.2 Dashboards

Dashboards consolidate key metrics into a single, often interactive, view for ongoing monitoring. Good dashboard design principles: prioritize the most important metrics visually (top-left, largest), avoid clutter, use consistent color coding, provide drill-down/filter capability, and ensure the audience can act on what they see.

Interview Q&A: - *Q: What makes a dashboard “bad”?* A: Too many metrics without hierarchy, inconsistent or misleading scales, chart types that don’t fit the data, no clear takeaway, and lack of context (no benchmarks or trends).

7.3 Data Visualization Principles

Choose chart types based on the message: - **Trend over time** → line chart. - **Comparison across categories** → bar chart. - **Part-to-whole** → stacked bar or (sparingly) pie chart. - **Relationship/correlation** → scatter plot. - **Distribution** → histogram or boxplot.

Best practices: avoid 3D effects and unnecessary chart junk, start bar charts at zero, use color purposefully (highlight, don’t decorate), and label directly rather than relying only on legends.

Interview Q&A: - *Q: Why are pie charts often discouraged?* A: Humans are poor at comparing angles/areas precisely; bar charts usually communicate the same part-to-whole comparison more accurately, especially with more than 3-4 categories.

7.4 Storytelling with Data

Data storytelling combines data, visuals, and narrative to drive a decision. Structure: start with the business context/question, present the key insight (not every finding), use visuals to support — not repeat — the narrative, and end with a clear, actionable recommendation.

Interview Q&A: - *Q: How do you tailor a data presentation to a non-technical stakeholder?* A: Lead with the takeaway and business impact rather than methodology, minimize jargon, use simple and clean visuals, and be ready to go deeper only if asked. - *Q: Describe a time you used data to influence a decision.* A: Strong answers follow a structure: context/problem → analysis approach → key insight → recommendation → measurable outcome. Interviewers are listening for business impact, not just technical execution.

Part 8: Common Interview Mistakes and Best Practices

8.1 Common Mistakes Candidates Make

1. **Jumping to a solution before clarifying the question.** Candidates often start writing SQL or opening Excel the moment a prompt is given. Always restate the

business problem first (“So we’re trying to understand X, using Y data, over Z timeframe — is that right?”) to confirm alignment before investing time in the wrong direction.

2. **Ignoring data quality.** Presenting results without mentioning missing values, duplicates, or outliers checked signals a lack of rigor. Even a brief caveat (“I noticed ~3% missing values in this field, which I excluded”) demonstrates care.
3. **Overcomplicating SQL/Excel solutions.** Using a nested subquery, multiple helper columns, or a convoluted VLOOKUP chain when a simple JOIN, CTE, or window function would do suggests unfamiliarity with more efficient tools — interviewers often ask “is there a simpler way?” as a direct probe.
4. **Confusing correlation with causation.** Claiming a metric “caused” an outcome based purely on a correlated trend, without considering confounders or experimental evidence, is a major red flag in analytics interviews.
5. **Not tying analysis back to business impact.** Reciting numbers without a “so what” — e.g., listing statistics without stating what decision they support — leaves interviewers unconvinced of business judgment.
6. **Weak handling of ambiguous case/take-home questions.** Not stating assumptions explicitly (e.g., “I assumed ‘active user’ means logged in within 30 days”) makes it hard for an interviewer to evaluate reasoning, and can make otherwise-correct answers look arbitrary.
7. **Poor whiteboard/live-SQL communication.** Coding silently instead of narrating logic prevents the interviewer from following (and crediting) your reasoning, even if you eventually reach the right answer.
8. **Ignoring edge cases.** Forgetting NULLs in JOINS, divide-by-zero risk in ratio metrics, or duplicate rows inflating SUM/COUNT results are common technical slip-ups that experienced interviewers specifically probe for.
9. **Overfitting the answer to sound impressive.** Using unnecessarily advanced terminology or techniques (e.g., invoking machine learning for a question answerable with a simple average) can come across as poor judgment about the right tool for the job.
10. **Not asking about data availability before proposing a plan.** Proposing an analysis plan that assumes data no one has confirmed exists (e.g., “we’ll segment by intent,” when intent isn’t tracked) undermines credibility — check feasibility before committing to an approach.

8.2 Best Practices for Interviews

- **Think aloud** during technical exercises — interviewers assess reasoning, not just the final answer.

- **Clarify ambiguous requirements** before diving into a case study (e.g., “What time period should I use?”).
- **State assumptions explicitly** when data or requirements are incomplete.
- **Validate results** with sanity checks (e.g., row counts before/after a JOIN, totals that reconcile).
- **Connect technical answers to business value** — every metric or query should tie back to a decision.
- **Know your resume cold** — be ready to explain the “why,” not just the “what,” of past projects.
- **Practice explaining trade-offs** — e.g., normalization vs. denormalization, mean vs. median, VLOOKUP vs. XLOOKUP — interviewers love trade-off reasoning.
- **Prepare STAR-format stories** (Situation, Task, Action, Result) for behavioral questions about data-driven decisions.

8.2b The STAR Framework for Behavioral Questions

STAR structures a behavioral answer into four parts:

- **Situation** — brief context: the team, project, or business problem.
- **Task** — the specific responsibility or goal you owned.
- **Action** — the concrete steps you personally took (use “I,” not just “we”).
- **Result** — the measurable outcome, ideally quantified (e.g., “reduced churn by 8%,” “saved 5 hours/week”).

Interview Q&A: - *Q: Why do interviewers favor the STAR format?* A: It forces a candidate to demonstrate ownership and impact rather than vague generalities, and gives the interviewer a clear, comparable structure to evaluate against other candidates. - *Q: What’s a common weakness in STAR answers?* A: Spending too long on Situation/Task and rushing the Result, or giving a result with no number attached — interviewers remember specific, quantified outcomes far more than general descriptions.

8.3 Sample Behavioral Interview Questions

- “Tell me about a time your analysis was wrong. What did you learn?”
- “Describe a project where the data didn’t support the expected hypothesis.”
- “How do you handle disagreement with a stakeholder about your findings?”
- “Tell me about a time you had to simplify a complex analysis for a non-technical audience.”

8.4 Sample Technical Interview Prompts

- “Write a SQL query to find the second-highest salary in each department.”
 - “Given this dataset, calculate month-over-month retention.”
 - “How would you detect and handle outliers in this sales dataset?”
 - “Design a dashboard for a product manager tracking user engagement.”
 - “Explain how you’d set up an A/B test for a new checkout flow.”
-

Appendix A: Quick-Reference Cheat Sheets

A.1 SQL Clause Execution Order

FROM → JOIN → WHERE → GROUP BY → HAVING → SELECT → DISTINCT → ORDER BY → LIMIT/OFFSET

A.2 Excel Lookup Function Comparison

Function	Direction	Match Type Default	Robust to Inserted Columns
VLOOKUP	Left-to-right only	Approximate	No
HLOOKUP	Top-to-bottom only	Approximate	No
INDEX/MATCH	Any direction	Exact (with 0)	Yes
XLOOKUP	Any direction	Exact	Yes

A.3 SQL JOIN Types at a Glance

JOIN Type	Returns
INNER JOIN	Only matching rows in both tables
LEFT JOIN	All left rows + matched right rows (NULL if unmatched)
RIGHT JOIN	All right rows + matched left rows (NULL if unmatched)
FULL OUTER JOIN	All rows from both, matched where possible
CROSS JOIN	Cartesian product (every combination)
SELF JOIN	Table joined to itself

A.4 Central Tendency and Spread Cheat Sheet

Statistic	Formula / Definition	Sensitive to Outliers?
Mean	Sum of values / count	Yes
Median	Middle value when sorted	No
Mode	Most frequent value	No
Variance	Average squared deviation from mean	Yes
Standard Deviation	$\sqrt{\text{Variance}}$	Yes
IQR	$Q3 - Q1$	No

A.5 Statistical Test Selection Guide

Goal	Test
Compare mean of a sample to a known value	One-sample t-test
Compare means of two independent groups	Two-sample (independent) t-test
Compare means of paired/before-after data	Paired t-test
Compare means across 3+ groups	ANOVA
Test association between two categorical variables	Chi-square test
Compare a proportion (e.g., conversion rate) between two groups	Two-proportion z-test

A.6 Chart Type Selection Guide

Goal	Recommended Chart
Trend over time	Line chart
Compare categories	Bar chart
Part-to-whole	Stacked bar (or pie for ≤ 3 categories)
Relationship between two numeric variables	Scatter plot
Distribution of a single variable	Histogram or boxplot
Comparing distributions across groups	Boxplot or violin plot
Geographic data	Map / choropleth

A.7 Pandas vs. SQL Equivalents

SQL	Pandas
SELECT col FROM df	df["col"]
WHERE condition	df[df["col"] > x]
GROUP BY col	df.groupby("col")
ORDER BY col	df.sort_values("col")
JOIN ... ON	df.merge(other, on="key", how="inner")
DISTINCT	df.drop_duplicates()
LIMIT n	df.head(n)

Appendix B: Glossary of Key Terms

A/B Test — A controlled experiment comparing two variants to measure impact on a metric. **Aggregate Function** — A function (SUM, AVG, COUNT, MIN, MAX) that summarizes multiple rows into one value. **Cardinality** — The uniqueness of values in a column, or the type of relationship (one-to-one, one-to-many, many-to-many) between tables. **CLT (Central Limit Theorem)** — Sample means approach a normal distribution as sample size increases, regardless of the population's distribution. **Confidence Interval** — A range of plausible values for a population parameter at a given confidence level. **Confounding Variable** — An unmeasured variable that influences both the independent and dependent variable, creating a spurious association. **CTE (Common Table Expression)** — A named temporary result set defined with WITH in SQL. **Data Grain** — What a single row in a table represents. **DAX** — Data Analysis Expressions, the formula language used in Power BI. **Denormalization** — Intentionally introducing redundancy to improve read performance. **ETL** — Extract, Transform, Load: the process of moving and reshaping data from source systems into an analysis-ready format. **Fact Table** — A table holding quantitative, transactional data in a star schema. **Foreign Key** — A column referencing a primary key in another table. **Index** — A data structure that speeds up lookups on specific columns. **KPI** — Key Performance Indicator: a measurable value tied to a business objective. **LOD Expression** — A Tableau calculation controlling the level of aggregation independent of the view. **Materialized View** — A view whose result set is physically stored and periodically refreshed. **Normalization** — Organizing relational data to reduce redundancy and improve integrity. **Null Hypothesis (H0)** — The default assumption of no effect or no difference in a hypothesis test. **Outlier** — A data point significantly different from the rest of the distribution. **p-value** — The probability of observing a

result as extreme as the sample, assuming the null hypothesis is true. **Primary Key** — A column (or set of columns) that uniquely identifies each row in a table. **Star Schema** — A data warehouse design with a central fact table connected to surrounding dimension tables. **Type I Error** — Rejecting a true null hypothesis (false positive). **Type II Error** — Failing to reject a false null hypothesis (false negative). **View** — A saved, virtual table defined by a SQL query. **Window Function** — A SQL function that computes a value across a set of related rows without collapsing them into one row. **Cohort Analysis** — Grouping users by a shared starting point (e.g., signup week) to track behavior over time. **Data Grain** — See above; the level of detail one row in a table represents. **ELT vs. ETL** — ELT loads raw data first and transforms it inside the warehouse; ETL transforms data before loading it, typically outside the warehouse. **Feature (in analytics context)** — An individual measurable attribute used as an input to analysis or modeling. **Funnel Analysis** — Measuring conversion rates between sequential steps of a process to identify drop-off points. **Idempotency** — A property of a data pipeline where re-running it produces the same result without duplicating data. **MECE** — Mutually Exclusive, Collectively Exhaustive; a framework for structuring categories so they don't overlap and fully cover all cases. **Retention Rate** — The percentage of users who return after a given period. **Sampling Bias** — Systematic error introduced when a sample isn't representative of the population it's drawn from. **Segmentation** — Dividing data into meaningful subgroups (e.g., by geography, plan tier, or acquisition channel) for more granular analysis. **Skewness** — A measure of the asymmetry of a distribution; positive skew has a long right tail, negative skew a long left tail. **Time Series** — Data indexed in time order, often analyzed for trend, seasonality, and anomalies.

Appendix C: Rapid-Fire Interview Question Bank

1. *What's the difference between WHERE and HAVING?* — WHERE filters rows before aggregation; HAVING filters groups after aggregation.
2. *What does GROUP BY do?* — Groups rows sharing a value so aggregate functions can summarize each group.
3. *What's the difference between COUNT() and COUNT(column)?** — COUNT(*) counts all rows including NULLs; COUNT(column) counts only non-NULL values in that column.
4. *What is a self join used for?* — Comparing rows within the same table, e.g., matching employees to their managers.
5. *What's the difference between DELETE and TRUNCATE?* — DELETE removes

rows conditionally and can be rolled back (logged row-by-row); TRUNCATE removes all rows at once and is typically faster but not row-logged.

6. *What is a primary key vs. a unique key?* — A table can have only one primary key (which also disallows NULLs); it can have multiple unique keys, which do allow a single NULL in most databases.
7. *What does INDEX/MATCH do that VLOOKUP cannot?* — It can look up values to the left of the lookup column and remains stable if columns are inserted or deleted.
8. *What is a pivot table calculated field?* — A custom formula that computes a new value from existing fields within the pivot table.
9. *What's the difference between .csv and .xlsx?* — CSV is a plain-text, single-sheet format with no formatting/formulas; XLSX is a binary/XML format supporting multiple sheets, formulas, and formatting.
10. *What is Power Query used for?* — Importing, cleaning, and transforming data from multiple sources in a repeatable, refreshable way.
11. *What's the difference between a list and a NumPy array?* — NumPy arrays are fixed-type, contiguous in memory, and support fast vectorized math; Python lists are flexible but slower for numerical operations.
12. *How do you check for missing values in pandas?* — `df.isnull().sum()` shows the count of missing values per column.
13. *What does `df.describe()` return?* — Summary statistics (count, mean, std, min, quartiles, max) for numeric columns.
14. *What is the difference between mean and median for skewed data?* — Mean is pulled toward the tail of the skew; median remains a more robust measure of the “typical” value.
15. *What is standard error and how does it differ from standard deviation?* — Standard deviation measures spread within a sample; standard error measures the precision of the sample mean as an estimate of the population mean, and shrinks as sample size grows.
16. *What is statistical power?* — The probability of correctly detecting an effect when one truly exists (1 – probability of a Type II error).
17. *What's a common cause of Simpson's Paradox in A/B testing?* — A trend appears in aggregated data but reverses (or disappears) when the data is split by a confounding subgroup, often due to unequal group sizes across segments.
18. *What's the difference between a KPI and a metric?* — All KPIs are metrics, but not all metrics are KPIs; a KPI is a metric explicitly tied to a strategic goal and tracked for decision-making.
19. *What is churn rate?* — The percentage of customers (or revenue) lost over a

given period, often calculated as customers lost $\div$ customers at the start of the period.

20. *What is Customer Lifetime Value (CLV)?* — The total revenue a business expects from a customer over the entire relationship, often estimated as average purchase value $\times$ purchase frequency $\times$ customer lifespan.
21. *Why avoid dual-axis charts in most cases?* — They can visually mislead by implying a relationship between two variables through arbitrary axis scaling; used carefully, they can be appropriate for genuinely related metrics.
22. *What is data storytelling?* — Combining data, visuals, and narrative to lead an audience to a clear, actionable conclusion.
23. *How would you explain a p-value to a non-technical stakeholder?* — “It tells us how likely we’d see a result this strong just by chance if there were actually no real effect — a small p-value means the effect is probably real.”
24. *What’s the difference between an INNER JOIN and a subquery with WHERE IN?* — Both can filter on matches from a related table, but a JOIN can return columns from both tables, while WHERE IN only filters the outer query’s rows and cannot pull extra columns from the subquery.
25. *What’s a common reason a dashboard number doesn’t match a raw SQL query?* — Different filter context, a different definition of the same metric (e.g., “active user” defined differently), timezone mismatches, or caching/refresh lag.
26. *What is the difference between an inner join and a semi-join?* — An inner join returns matching rows with columns from both tables; a semi-join (e.g., WHERE EXISTS) filters the left table to only rows with a match, but returns no columns from the right table.
27. *What is data leakage in analysis or modeling?* — When information that wouldn’t be available at prediction/decision time is accidentally included, making results look better than they truly are.
28. *What’s the difference between a histogram and a bar chart?* — A histogram shows the distribution of a continuous numeric variable grouped into bins; a bar chart compares discrete categories.
29. *What does “statistically significant but not practically significant” mean?* — The observed effect is unlikely due to chance, but is too small in magnitude to justify business action or cost.
30. *What’s the difference between ROLLUP/CUBE and GROUP BY?* — ROLLUP/CUBE generate subtotal and grand-total rows automatically across grouping levels, whereas plain GROUP BY returns only the detail-level groups.
31. *Why might you choose a median-based metric for a highly skewed KPI like “time to purchase”?* — Because a few extreme outliers would distort the mean, while

the median better reflects the typical customer experience.

32. *What is the difference between a wide format and a long (tidy) format dataset?* — Wide format has one row per subject with multiple columns per variable/time period; long format has one row per subject-variable(-time) combination — most plotting and modeling libraries expect long/tidy format.
33. *What's the purpose of a data dictionary?* — Documentation describing each field's meaning, type, and valid values, ensuring consistent interpretation of data across a team.
34. *What is survivorship bias, and how might it appear in analysis?* — Drawing conclusions only from the subset of “surviving” cases while ignoring those that dropped out, which can bias findings toward overly positive conclusions.
35. *What's the difference between a report and a dashboard?* — A report is typically a static, point-in-time summary; a dashboard is a live, interactive tool for ongoing monitoring, usually refreshed automatically.

Excel Keyboard Shortcuts Reference

Shortcut (Windows)	Action
Ctrl + Shift + L	Toggle filters
Ctrl + T	Create a table
Ctrl + Shift + Arrow	Select to end of data range
Alt + =	AutoSum
Ctrl + Shift + %	Apply percentage format
F4	Toggle absolute/relative reference
Ctrl + ;	Insert current date
Alt + F1	Insert quick chart

Appendix D: Worked SQL Practice Problems

Problem 1: Find the second-highest salary in each department.

```
SELECT department, salary
FROM (
    SELECT department, salary,
           DENSE_RANK() OVER (PARTITION BY department ORDER BY salary DESC) AS rnk
    FROM employees
```

```
) ranked  
WHERE rnk = 2;
```

Why DENSE_RANK: Using DENSE_RANK (rather than ROW_NUMBER) correctly handles ties — if two employees share the highest salary, the “second-highest” is the next distinct value, not simply the third row.

Problem 2: Find customers who placed orders in every month of 2026.

```
SELECT customer_id  
FROM orders  
WHERE EXTRACT(YEAR FROM order_date) = 2026  
GROUP BY customer_id  
HAVING COUNT(DISTINCT EXTRACT(MONTH FROM order_date)) = 12;
```

Why this works: Counting distinct months per customer and comparing to 12 confirms full-year coverage without needing to check each month individually.

Problem 3: Calculate running total of daily sales.

```
SELECT sale_date, daily_sales,  
       SUM(daily_sales) OVER (ORDER BY sale_date) AS running_total  
FROM daily_sales_summary;
```

Problem 4: Find duplicate rows based on a set of columns.

```
SELECT customer_id, email, COUNT(*)  
FROM customers  
GROUP BY customer_id, email  
HAVING COUNT(*) > 1;
```

Problem 5: Calculate the percentage each product contributes to total category sales.

```
SELECT product_id, category, sales,  
       sales / SUM(sales) OVER (PARTITION BY category) AS pct_of_category  
FROM product_sales;
```

Appendix E: Worked Python/Pandas Practice Problems

Problem 1: Find the top 3 products by revenue per category.

```
top3 = (
    df.sort_values("revenue", ascending=False)
      .groupby("category")
      .head(3)
)
```

Problem 2: Calculate month-over-month percentage change in revenue.

```
monthly = df.groupby("month")["revenue"].sum().reset_index()
monthly["mom_change"] = monthly["revenue"].pct_change()
```

Problem 3: Flag customers as churned if they haven't ordered in 90+ days.

```
latest_order = df.groupby("customer_id")["order_date"].max().reset_index()
latest_order["days_since_order"] = (pd.Timestamp("today") - latest_order["order_date"])
latest_order["churned"] = latest_order["days_since_order"] > 90
```

Problem 4: Clean a messy text column (extra spaces, inconsistent casing).

```
df["category"] = df["category"].str.strip().str.lower().str.replace(r"\s+", " ", regex=
```

Problem 5: Merge two datasets and identify unmatched rows (anti-join).

```
merged = df1.merge(df2, on="customer_id", how="left", indicator=True)
unmatched = merged[merged["_merge"] == "left_only"]
```

Appendix F: Sample Case Study Walkthrough

Prompt: “Weekly active users (WAU) dropped 12% last week. Walk me through how you’d investigate this.”

Step 1 — Clarify the metric and scope. Confirm the exact definition of WAU, the comparison period (week-over-week vs. same week last year), and whether the drop is company-wide or specific to a platform, region, or user segment.

Step 2 — Rule out data issues first. Check for tracking/logging errors, pipeline failures, or a change in how “active” is defined before assuming a genuine behavioral drop — a large sudden change is often a data issue, not a real trend.

Step 3 — Segment the drop. Break WAU down by platform (iOS/Android/Web), acquisition channel, geography, and new vs. returning users. This isolates whether the decline is broad-based or concentrated in one segment.

Step 4 — Check for known events. Cross-reference the timeline with product releases, marketing campaign changes, pricing changes, outages, or seasonality (e.g., holidays).

Step 5 — Quantify and prioritize. Rank segments by their contribution to the overall drop (e.g., “70% of the decline came from Android users in the EU”) rather than presenting every cut of the data equally.

Step 6 — Recommend next steps. Propose a specific action: further root-cause analysis, a fix to a suspected bug, or a monitoring plan — tied to expected impact.

What interviewers are evaluating: structured thinking, prioritization under ambiguity, and the ability to move from “the number changed” to “here’s the likely cause and what to do about it” — not just data manipulation skill.

Appendix G: Portfolio Project Ideas for Data Analysts

Building 2-3 well-documented projects is one of the highest-leverage ways to prepare for interviews, since they generate concrete stories for behavioral questions and demonstrate end-to-end thinking.

1. **Sales/Revenue Dashboard** — Use a public retail dataset to build a Power BI or Tableau dashboard tracking revenue, growth, and top products/regions, with a written summary of key insights.
2. **A/B Test Analysis** — Simulate or find an experiment dataset, calculate conversion rates by group, run a significance test, and write a recommendation memo including sample size and confidence interval reasoning.
3. **Customer Churn Exploration** — Use Python/pandas to clean a subscription dataset, calculate churn rate by segment, and visualize the biggest churn drivers.
4. **SQL Portfolio Queries** — Solve and document 10-15 progressively harder SQL problems (joins, window functions, CTEs) against a public database (e.g., a sample e-commerce schema), explaining the business question each query answers.
5. **Cohort Retention Analysis** — Build a cohort retention heatmap in Excel or Python showing how retention evolves by signup month.
6. **End-to-End Case Study Write-Up** — Pick a real or hypothetical business question, walk through the full data analysis process (Part 1 of this guide), and present findings as a short slide deck — useful for practicing storytelling with data.

Interview Q&A: - *Q: What should I highlight when discussing a portfolio project in an interview?* A: The business question you set out to answer, key decisions/trade-offs you made (not just the final chart), and the concrete insight or recommendation that resulted — technical steps matter less than the reasoning behind them.

Appendix H: 30-Day Interview Preparation Checklist

Week 1 — Foundations: Review SQL fundamentals (SELECT, WHERE, GROUP BY, JOINS); refresh Excel core functions and Pivot Tables; skim statistics fundamentals (mean/median/mode, standard deviation).

Week 2 — Intermediate Skills: Practice SQL window functions, CTEs, and subqueries daily; build one pandas data-cleaning exercise; learn XLOOKUP/Power Query if unfamiliar.

Week 3 — Applied Practice: Complete 2-3 end-to-end case studies (define question → analyze → visualize → recommend); practice explaining hypothesis testing and A/B testing out loud; build or polish one dashboard.

Week 4 — Interview Simulation: Do timed mock SQL/Excel exercises; rehearse 5-6 STAR-format behavioral stories; practice a 2-minute “walk me through a project” pitch; review this guide’s Rapid-Fire Question Bank (Appendix C) end to end.

Interview Q&A: - *Q: How much time should I spend on technical vs. behavioral prep?* A: Most analyst interviews weight both meaningfully — a common guideline is roughly balanced prep time, since strong technical skills without clear communication (and vice versa) both commonly cause otherwise-qualified candidates to fall short.

Appendix I: Skill Expectations by Experience Level

Interview bars vary by seniority. This section helps calibrate expectations.

Entry-Level / Junior Analyst: - Comfortable with core SQL (SELECT, WHERE, GROUP BY, basic JOINS) and Excel (formulas, Pivot Tables). - Can perform basic data cleaning and produce simple visualizations. - Communicates findings clearly but may need guidance on scoping ambiguous questions.

Mid-Level Analyst: - Fluent in window functions, CTEs, subqueries, and query optimization basics. - Independently scopes ambiguous business questions and

chooses appropriate statistical methods. - Builds and maintains dashboards in Power BI/Tableau; comfortable with Python/pandas for cleaning and analysis. - Can design and interpret A/B tests, including basic power/sample-size reasoning.

Senior Analyst / Analytics Lead: - Designs data models (star schemas) and defines metric standards across teams. - Mentors junior analysts and reviews their analysis for rigor and business relevance. - Partners directly with leadership to shape strategy using data, including trade-off discussions on measurement and experimentation design. - Anticipates data quality and scalability issues before they affect stakeholders.

Interview Q&A: - *Q: What separates a senior analyst from a mid-level analyst in interviews?* A: Seniority is usually demonstrated through judgment — knowing which question is worth answering, anticipating stakeholder needs, and reasoning about trade-offs — rather than only technical execution speed.

Closing Notes for the Interview Coach

When answering candidate questions, prioritize concise, structured explanations followed by a concrete example or interview Q&A pairing. Where relevant, connect technical concepts (SQL, Excel, Python, statistics) back to business reasoning and communication — this is consistently what differentiates strong Data Analyst candidates in interviews.